
DT2119 - Speech and speaker recognition - Lab 1 - Feature extraction

Table of Contents

.....	1
1.1 Enframe the audio signal	1
1.2 Pre-emphasys	3
1.3 Hamming window	4
1.4 Fast Fourier Transform	6
1.5 Mel filterbank log spectrum	8
1.6 Cosine Transform	10
1.7 Liftering	11
1.8 Calculation of the MFCCs for each utterance	17
2. Study the correlation between utterances	18
3. Distances	19
Linkage clustering	21
Train	22

Mel Frequency Cepstral Coefficients (MFCCs) are coefficients used in Speech recognition based on human auditive perception. These coefficients come from the need, in the field of autonomous audio recognition, to extract the main features of an audio signal while discarding all the irrelevant features that will make the recognition harder to achieve (background noise, emotion ...)

The sounds generated by a human depend on the shape of the vocal tract, position of the tongue, teeth, lips etc. Then, if we can determine this shape accurately, the phoneme that was produced will be easy to identify. The main objective of the MFCC is to accurately represent the envelope of the short time power spectrum in order to establish the shape of the vocal tract. Consisting then of one of the most important features in Speech recognition

First we need to prepare the workspace and load the data set

```
clear all;
close all;
load('tidigits.mat')
load('example.mat')
```

1.1 Enframe the audio signal

The purpose of this first step is to cut the original sample in many smaller ones. In this case, we want windowframes of 20 ms with a shift between windows of 10 ms. As the shift is smaller than the window, the computed frames will have shared parts between each other

Compute the number of samples per frame and shift

As our signal has been sampled, we don't have a continuous signal but a discrete one. This is why we are interested in knowing the number of samples per window rather than the time.

The sampling rate is $S = 20$ kHz. Then the period is $T = 1/S$. We have that:

Length of the window: $0.02T = 400$ samples Length of the shift: $0.01T = 200$ samples

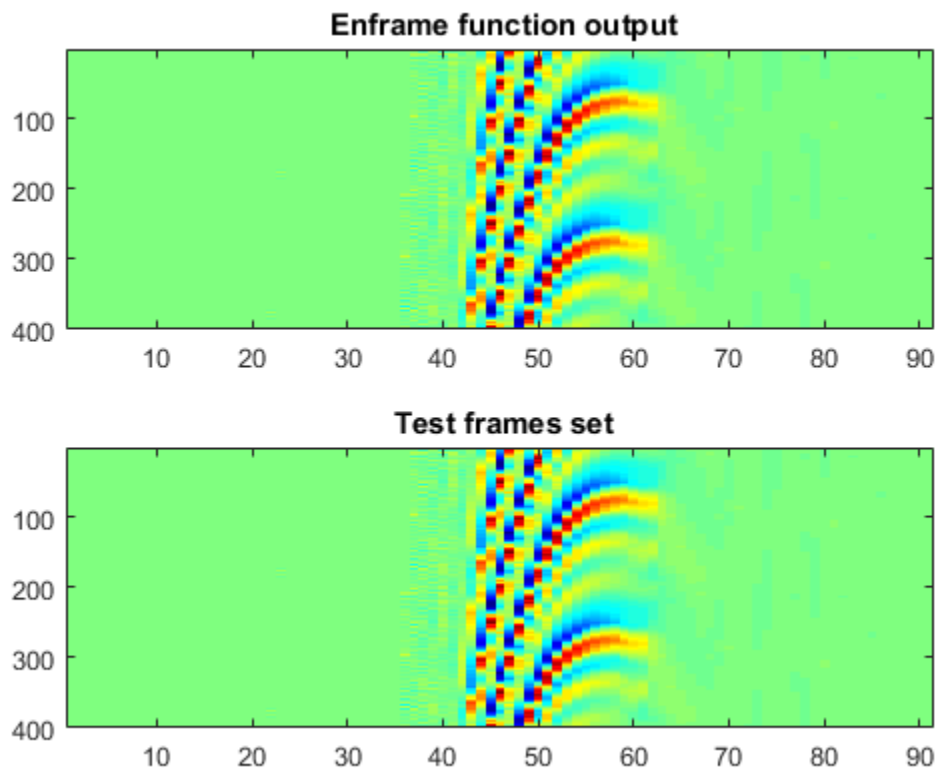
```
winlen = 400;  
winshift = 200;
```

Now is time to check the performance of the written function. To do so, the function is evaluated with the example's structure data set, and compared to the frames given in the example's structure frames. The output generated must be identical to the frames stored in the data structure in example.

```
samples = example{1,1}.samples;  
test = example{1,1}.frames';  
frames = enframe(samples,winlen,winshift);  
% As the image have to be plotted in with the same orientation as in  
% the  
% test data then  
frames = frames(:,1:400)';
```

The comparison is shown in the following figure.

```
subplot(2,1,1)  
imagesc(frames)  
colormap jet  
title(' Enframe function output')  
subplot(2,1,2)  
imagesc(test)  
colormap jet  
title(' Test frames set')
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test - frames;
Enframe_error = sum(sum(subs))
```

```
Enframe_error =
```

```
0
```

1.2 Pre-emphasys

The main objective of this function is to compensate the 6dB/octave that are dropped due to the radiation at the lips. This function is:

$$y[n] = x[n] - ax[n-1]$$

Being the coefficients $A = 1$ and $B = [1 - a]$.

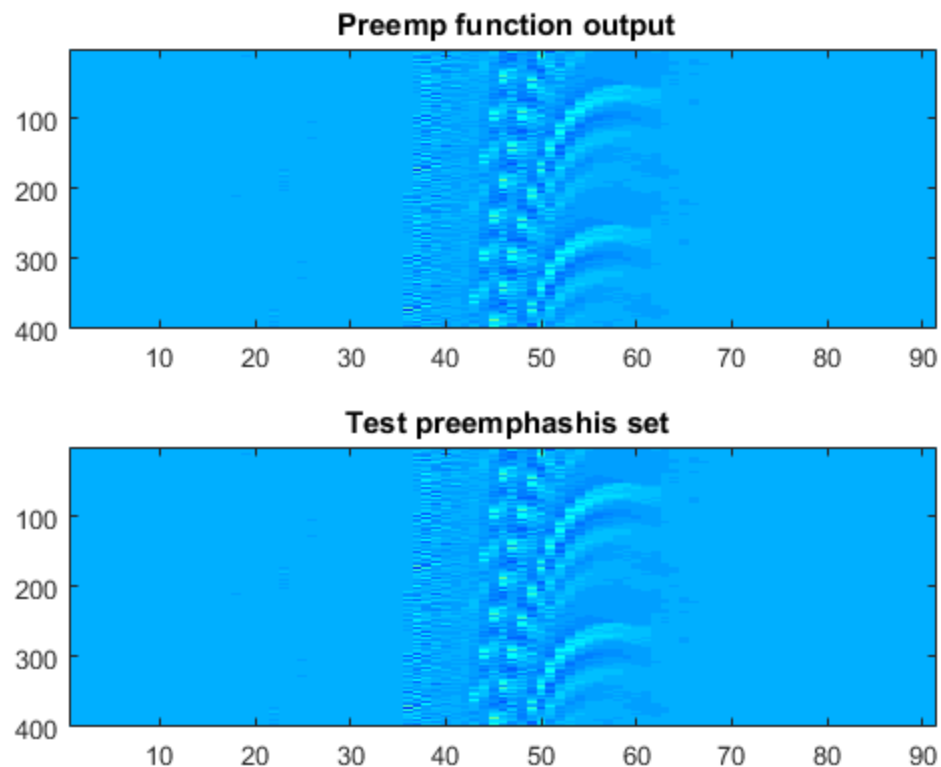
The purpose of this filter is to filter the lower frequencies, letting pass the higher ones. This is because the important features are in the higher frequencies, remaining the lower frequencies almost unchanged. Therefore, by doing the filtering we discard the similar parts of the signal analyzing only the higher frequencies, where the important features are.

We have now to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
a = 0.97;
preemph = preemp(frames,a);
test_preem = example{1,1}.preemph';
```

The comparison is shown in the following figure.

```
figure;
subplot(2,1,1)
imagesc(preemph)
colormap jet
title(' Preemp function output')
subplot(2,1,2)
imagesc(test_preem)
colormap jet
title(' Test preemphashis set')
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_preem - preemph;  
Preemphasis_error = sum(sum(subs))
```

Preemphasis_error =

0

1.3 Hamming window

The use of the Hamming window is justified as it emphasizes the center of the window, improving this way the values of the Fourier Transfer Function. The sidelobes are reduced, being the center lobe much more important than these side lobes.

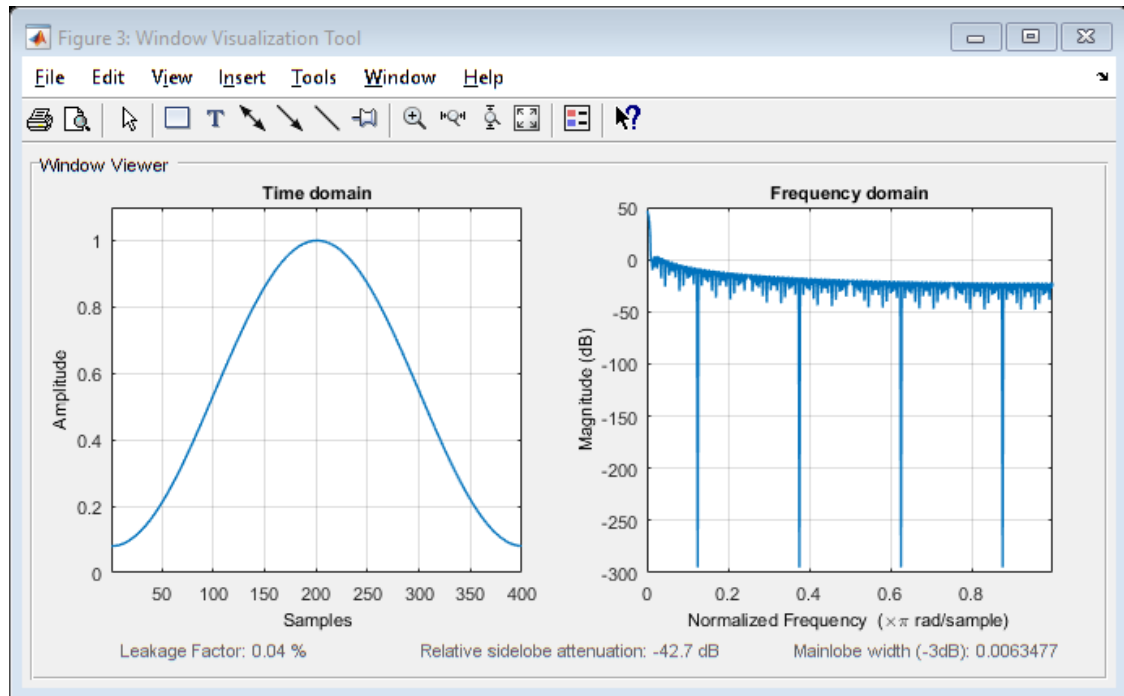
The Hamming window is used in order to reduce the discontinuities in the edges of the frames, focusing in the main(centered) frequencies.

The function windowing that applies the hamming window is implemented and used on the pre emphasized frames.

```
[windowed_frames, window] = windowing(preemph);  
test_windowed = example{1,1}.windowed';
```

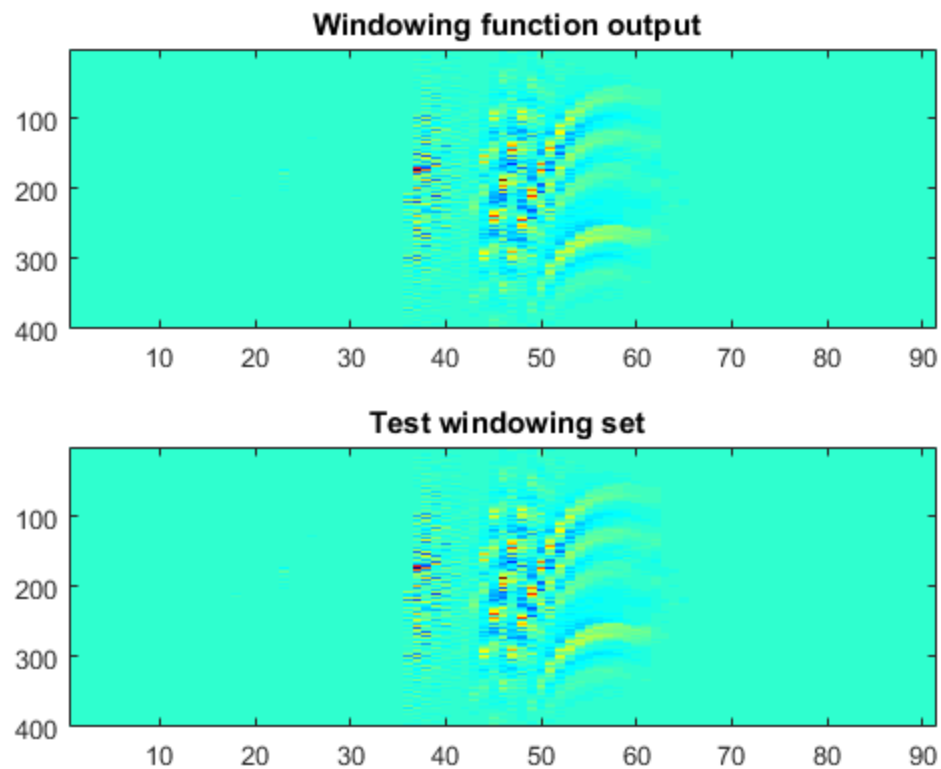
The shape of the Hamming window used is figure;

```
wvtool(window)
```



We now have to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
figure;  
subplot(2,1,1)  
imagesc(windowed_frames)  
colormap jet  
title(' Windowing function output')  
subplot(2,1,2)  
imagesc(test_windowed)  
colormap jet  
title(' Test windowing set')
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_windowed - windowed_frames;  
window_error = sum(sum(subs))
```

```
window_error =
```

```
-7.791701311932187e-13
```

We can observe that the error is not exactly zero but is sufficiently small enough. The insignificant difference is probably due to the use of Matlab instead of Python.

1.4 Fast Fourier Transform

The powerSpectrum function is developed in this section. This function performs the Fast Fourier Transform to the windowed frames and then applies a squared power to the modulus. The FFT samples length is 512.

As now we are in the frequency domain, it is interesting what the max frequency will be:

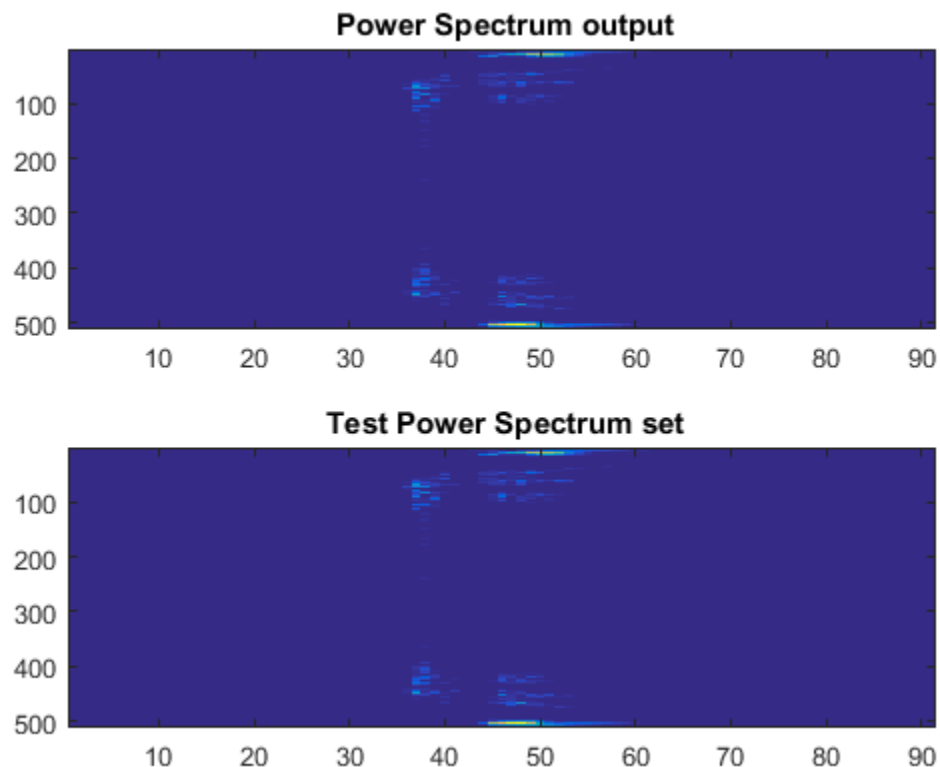
$$f_{max} = \frac{f_{samp}}{2} = \frac{20000}{2} = 10kHz$$

```
nfft = 512;
```

```
FFT_frame = powerSpectrum(windowed_frames,nfft);  
test_FFT = example{1,1}.spec';
```

We have now to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
figure;  
subplot(2,1,1)  
imagesc(FFT_frame)  
title(' Power Spectrum output')  
subplot(2,1,2)  
imagesc(test_FFT)  
title(' Test Power Spectrum set')
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_FFT - FFT_frame;  
FFT_error = sum(sum(subs))
```

FFT_error =

-2.497053070175601e-07

As expected, the error is not 0 as we get the errors from the previous step, but squared. Still the error is too small to be considered.

1.5 Mel filterbank log spectrum

A filter bank is a set of filters. In the case of the Mel filterbank, these filters are triangles, in which the first filter only keeps low frequencies while the next filters slowly shift to higher frequencies and the amplitude of the triangles decrease, and they get wider.

This decreased amplitude try to represent the human ear system, as the higher the frequencies, the more difficult discern differences. This is why the triangles get wider: we need a bigger filter as the resolution is less precise.

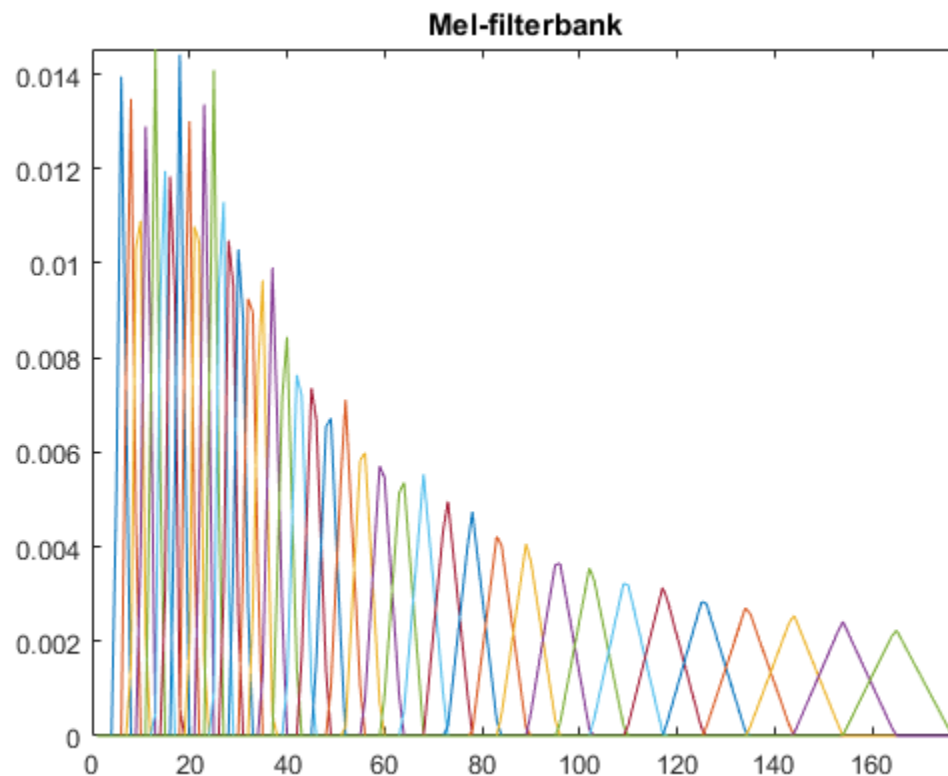
The sum of the energy in every filter is computed. We need to take the log of this energy due to the log property of sound: to double the sound we need 8 times more energy.

In this case the number of filters in the bank is 40: 13 linear filters and 27 logarithmic.

```
sampling_freq = 20000;  
[melSpec_frames, filterBank] =  
    logMelSpectrum(FFT_frame, sampling_freq);  
test_melSpec = example{1,1}.mspec';
```

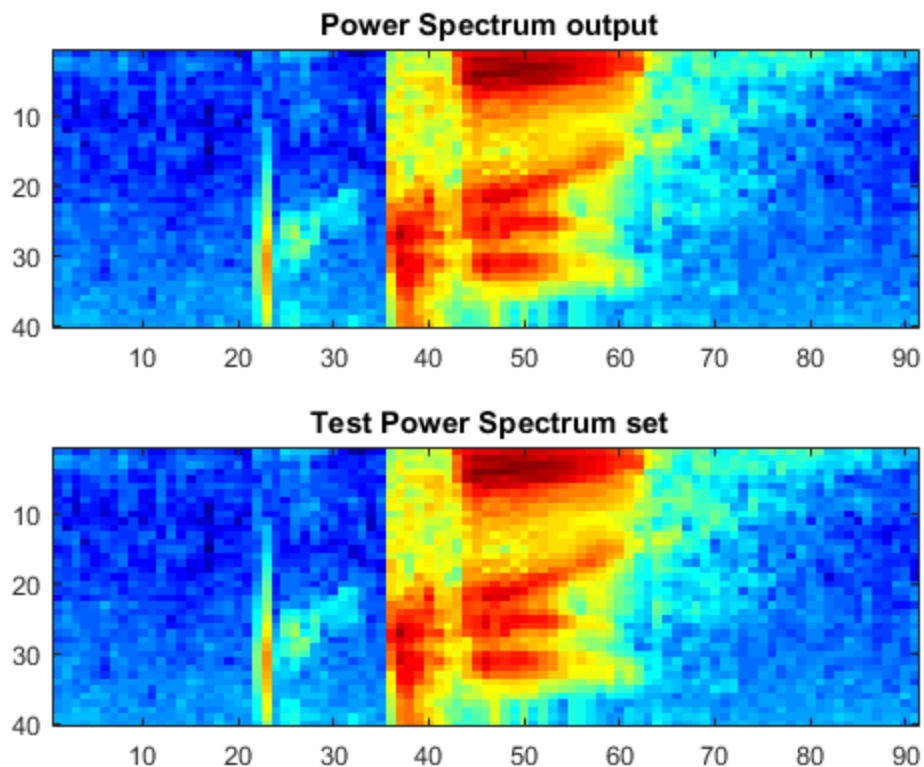
The Mel filter bank computed is the one that follows

```
figure('Name','Mel Frequency Filter Bank')  
plot(filterBank')  
axis([0, 178, 0, max(max(filterBank))])  
title('Mel-filterbank')
```



We have now to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
figure;  
subplot(2,1,1)  
imagesc(melSpec_frames)  
title(' Power Spectrum output')  
colormap jet  
subplot(2,1,2)  
imagesc(test_melSpec)  
title(' Test Power Spectrum set')  
colormap jet
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_melSpec - melSpec_frames;  
melSpec_error = sum(sum(subs))
```

```
melSpec_error =
```

```
3.582160609805207e-14
```

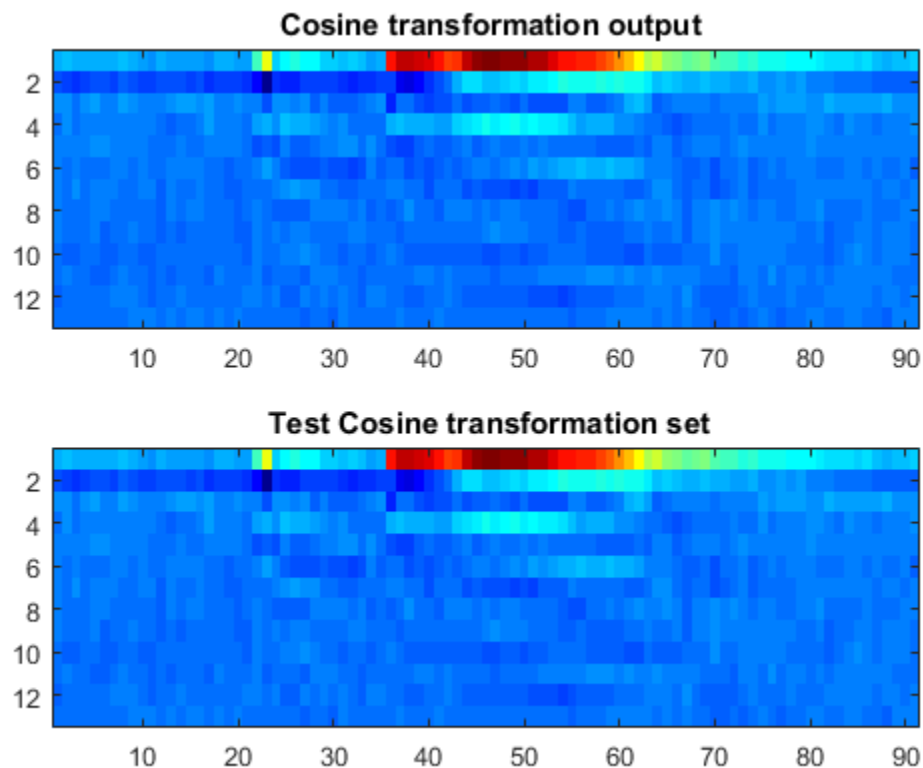
1.6 Cosine Transform

The Discrete Cosine Transform perform a similar operation than the Fourier transform: they both decompose a discrete-time vector in a sum of scaled-and-shifted basis functions. The main difference is that the DCT uses only cosines as the Basis function. We get then, for each window, 13 coefficients or cepstrums.

```
number_of_coefficients = 13;
[CT_frames] = cepstrum(melSpec_frames, number_of_coefficients);
test_CT = example{1,1}.mfcc';
```

We have now to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
figure;
subplot(2,1,1)
imagesc(CT_frames)
title(' Cosine transformation output')
colormap jet
subplot(2,1,2)
imagesc(test_CT)
title(' Test Cosine transformation set')
colormap jet
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_CT - CT_frames;
```

```
CT_error = sum(sum(subs))
```

```
CT_error =
```

```
4.555982327514307e-13
```

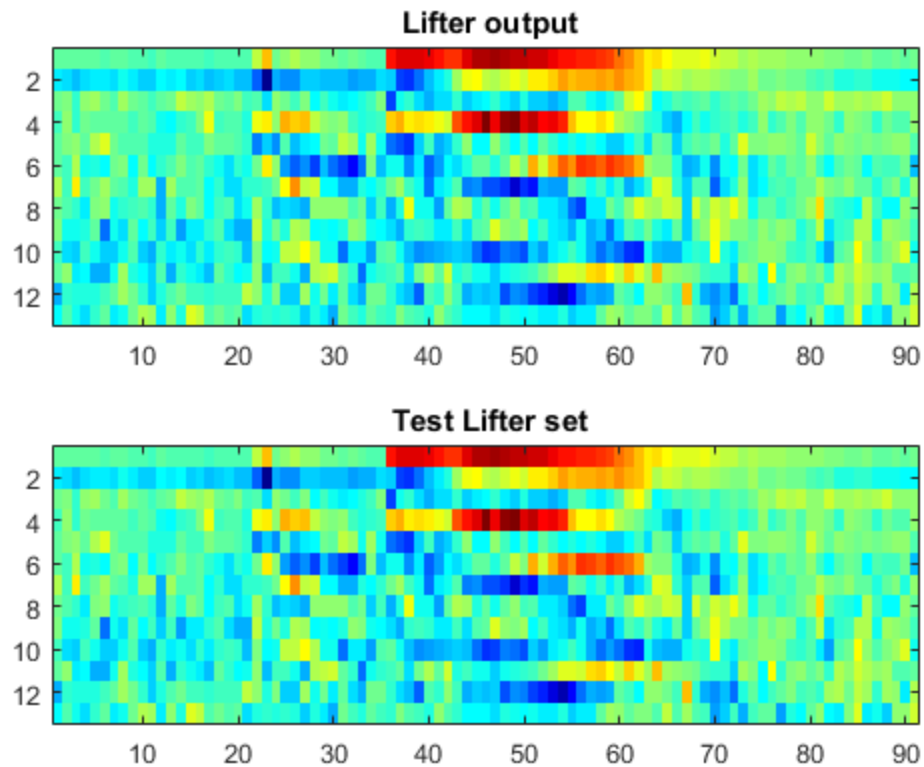
1.7 Liftering

This is the last step, which output are the final MFCCs for the current utterance. A lifter is a filter that operates on a cepstrum might be called a lifter. A low-pass lifter is similar to a low-pass filter in the frequency domain. This is done to even clean the signal more leaving only the important and different information.

```
[Lift_frames] = lifter_matlab(CT_frames);  
test_Lifter = example{1,1}.lmfcc';
```

We have now to check the performance of the written function. To do so, we have to compare the output of this function with the test data.

```
figure;  
subplot(2,1,1)  
imagesc(Lift_frames)  
title(' Lifter output')  
colormap jet  
subplot(2,1,2)  
imagesc(test_Lifter)  
title(' Test Lifter set')  
colormap jet
```



To check the performance numerically, we will subtract the test data from the output and get the total differences between both of them

```
subs = test_Lifter - Lift_frames;  
Lift_error = sum(sum(subs))
```

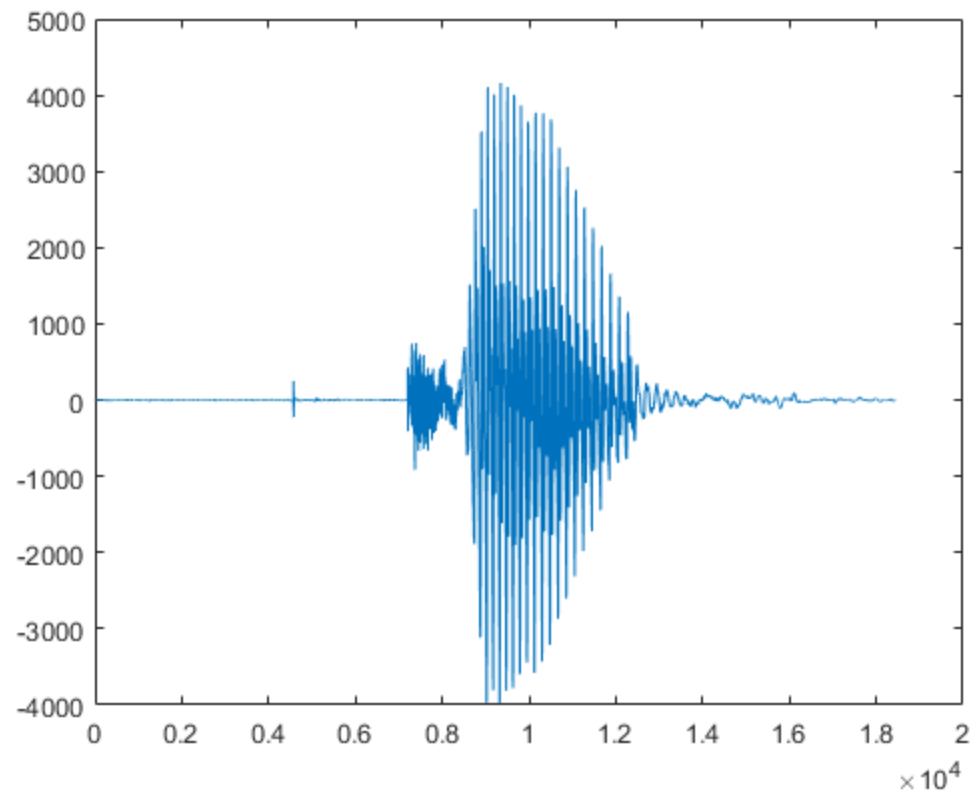
Lift_error =


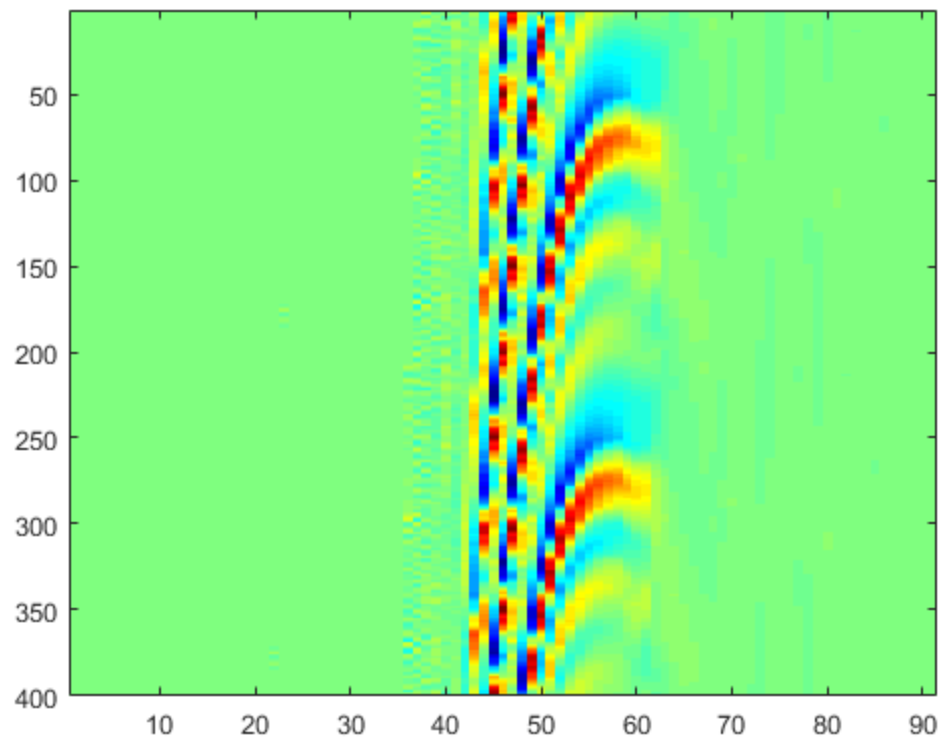
8.972197984569164e-13

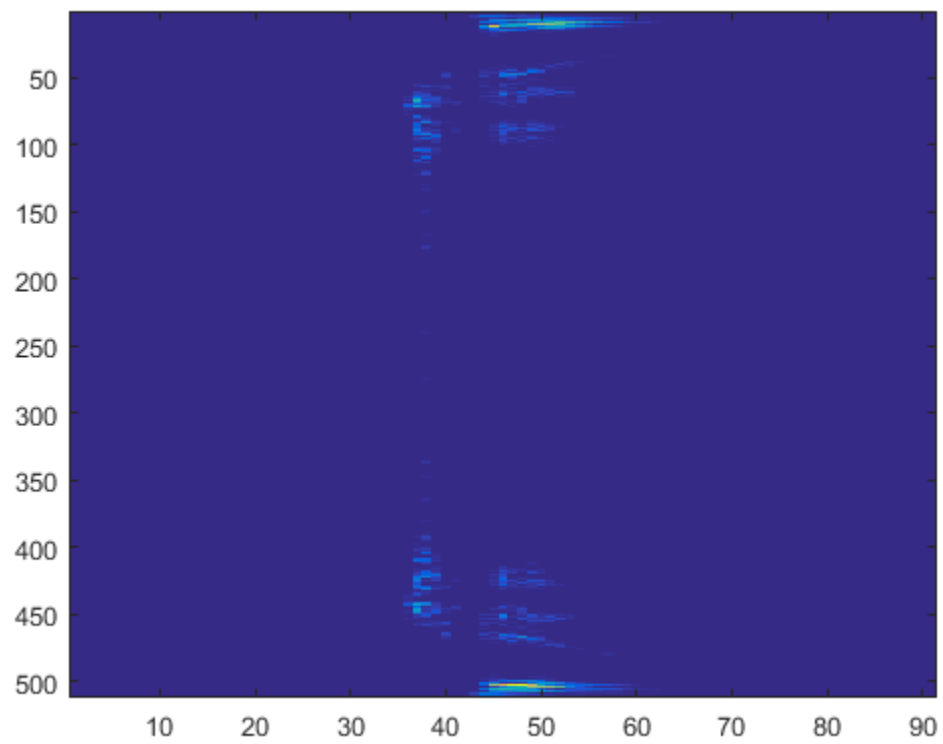
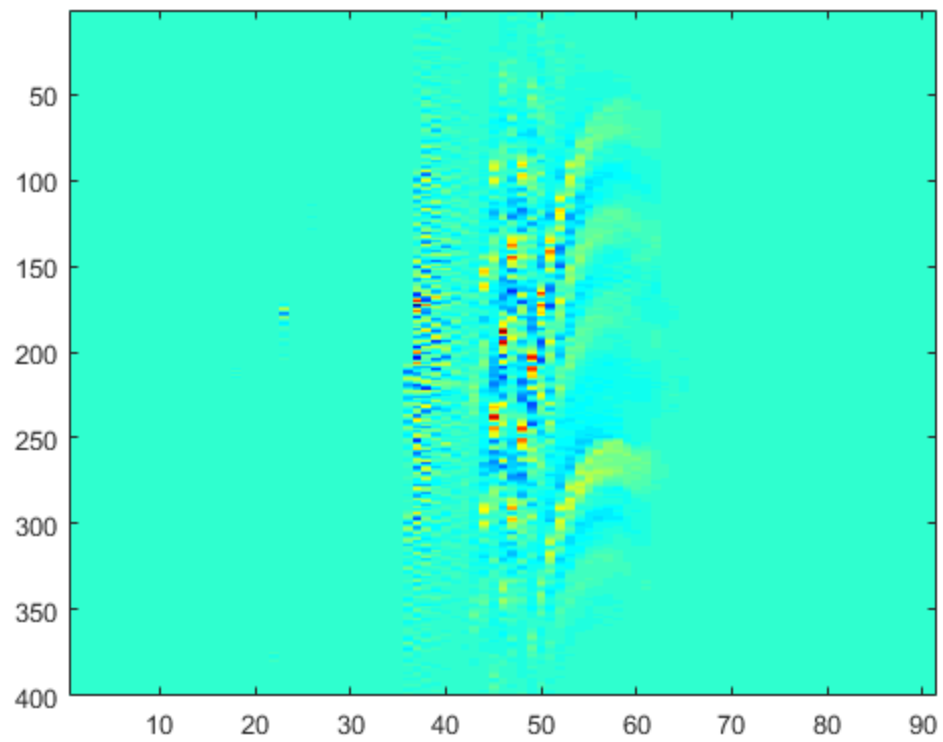
All the steps together in order can be represented by the following figure:


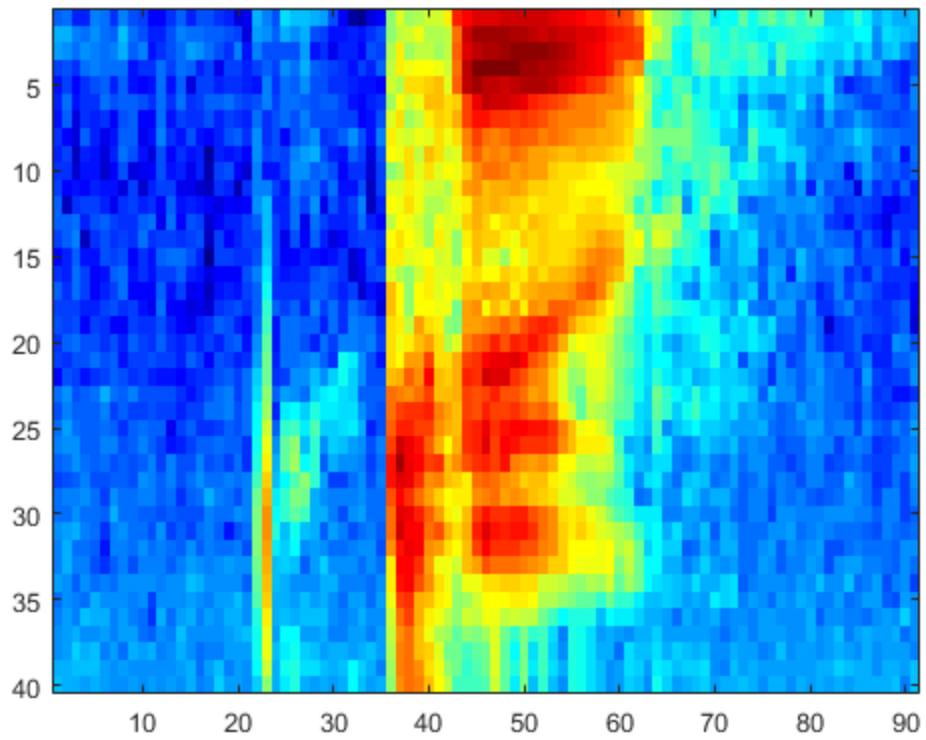
```
figure  
plot(samples)  
figure  
imagesc(frames);  
colormap(jet)  
figure  
imagesc(preemph)  
colormap(jet)  
figure  
imagesc(windowed_frames)  
colormap(jet)  
figure  
imagesc(FFT_frame)
```


```
colormap default
figure
imagesc(melSpec_frames)
colormap(jet)
figure
imagesc(CT_frames)
colormap(jet)
figure
imagesc(Lift_frames)
colormap(jet)
```











1.8 Calculation of the MFCCs for each utterance

For the further analysis of the data set, all the utterance are going to be transformed to the MFCCs space. Also, the output of each utterance will be concatenated in an unique matrix of features in order to be able to represent and gather all the features included in the data set. This matrix has dimensions of $N \times M$ being N the total number of frames in the data set and M the number of coefficients (13).

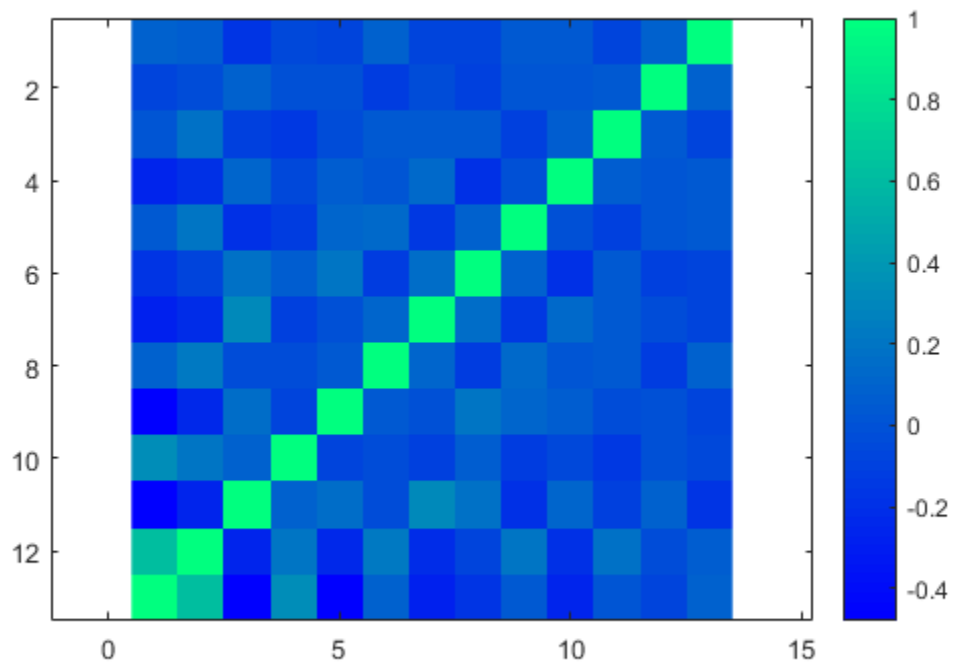
```
gender = [];  
speaker = [];  
digits = [];  
MFCCs = cell(1,44);  
mSpec = cell(1,44);  
MFCCs_straight = cell(1,44);  
mSpec_straight = cell(1,44);  
MFCCs_concatenate = [];  
mSpec_concatenate = [];  
for i = 1:size(tidigits,2)  
    [MFCCs_straight{i} mSpec_straight{i}] =  
    mfcc(tidigits{1,i}.samples);  
    MFCCs{i} = MFCCs_straight{i}';  
    MFCCs_concatenate = vertcat(MFCCs_concatenate,MFCCs{i});  
    mSpec{i} = mSpec_straight{i}';  
    mSpec_concatenate = vertcat(mSpec_concatenate,mSpec{i});  
    digits = [digits; tidigits{1,i}.digit,tidigits{1,i}.gender(1)];  
    gender = [gender; tidigits{1,i}.gender(1)];  
    speaker = [speaker; tidigits{1,i}.speaker];  
end
```

end

2. Study the correlation between utterances

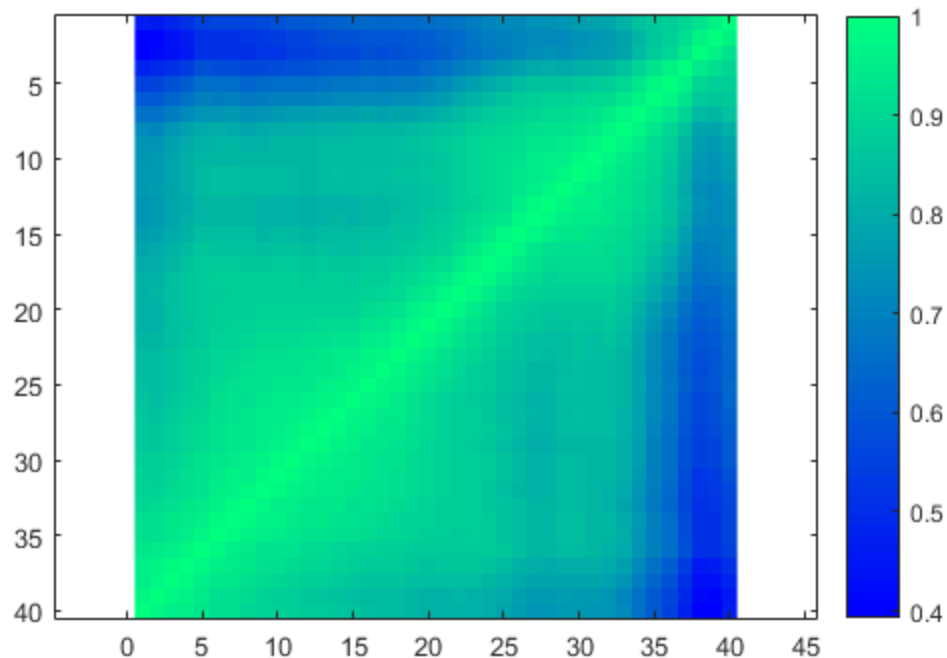
To study the correlation between the 13 coefficients, the correlation matrix of the entire feature matrix is calculated using the MFCCs coefficients.

```
figure
corMFCCs = corrcoef(MFCCs_concatenate);
imagesc(flipdim(corMFCCs ,1));           %# vertical flip)
axis equal
colormap winter
colorbar
```



And the Mel filterbank features

```
figure
corMspec = corrcoef(mSpec_concatenate);
imagesc(flipdim(corMspec ,1));           %# vertical flip)
axis equal
colormap winter
colorbar
```



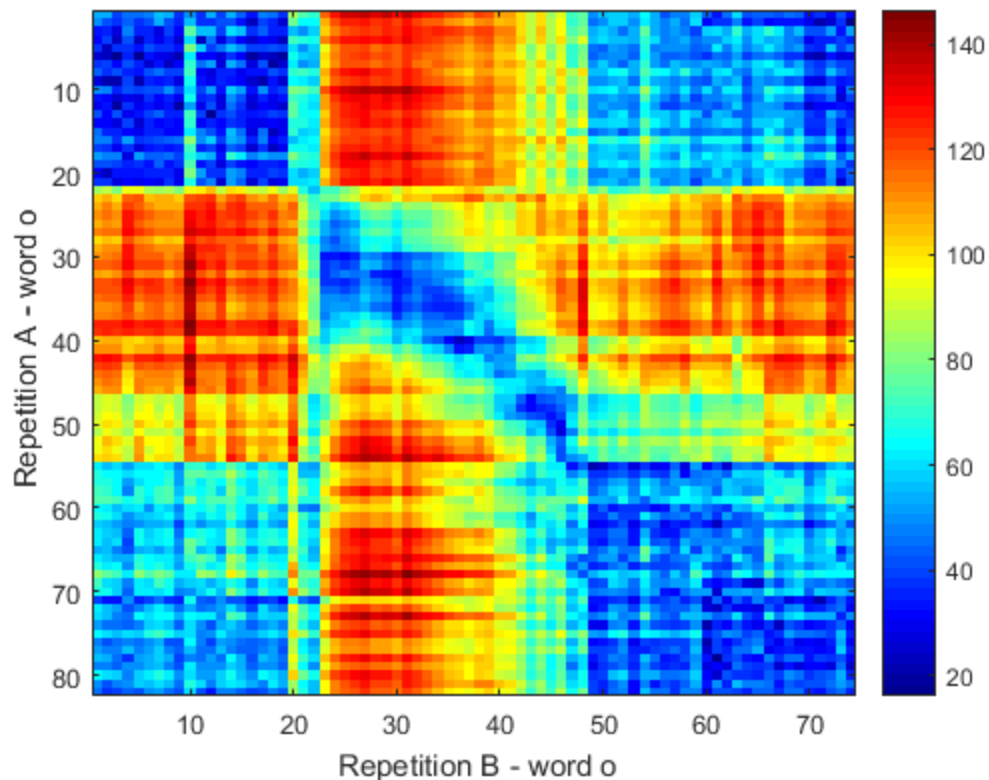
By comparing both results we can justify the steps that follow the application of the filterbank as all the correlations between coefficients have almost disappeared meaning this that all the coefficients are giving unique information. We can accept the diagonal covariance matrices assumption for Gaussian models as the MFCCs coefficients are almost not correlated at all.

3. Distances

The first step is that, given 2 utterances's MFCCs matrices of $N \times M$ each of the same word pronounce by the same speaker but in different repetitions, compute the Euclidean distance between them.

```
Rep_oa = mfcc(tidigits{1}.samples)';
Rep_ob = mfcc(tidigits{2}.samples)';

local_dist_o = localDistances(Rep_oa, Rep_ob);
figure
imagesc(local_dist_o)
xlabel(' Repetition B - word o' )
ylabel(' Repetition A - word o' )
colormap jet
colorbar
```



We can observe that in the first repetition the speaker said the word faster as the square that represent when the word is pronounced and therefore where the max distances are, is moved towards the upleft corner.

The last step is to apply the Dynamic Time Warping to compute the global distances between all the utterances's MFCCs. This step is done in order to study the relation between the digits, speakers and even genres and doing a later clustering.

```
D = zeros(size(tidigits,2),size(tidigits,2));
for i = 1:size(tidigits,2)
    for j = 1:size(tidigits,2)
        a = MFCCs{i};
        b = MFCCs{j};
        local_dist = localDistances(a, b);
        D(i,j) = dtw(local_dist);
    end
end
```

Once this Global distances matrix is computed, is time to plot it:

```
figure;
imagesc(D)
% colormap default(1000))
colormap(jet(60))
colorbar;
labels = tidigit2labels(tidigits);
set(gca,'XTick',1:44);
set(gca,'XTickLabel',labels);
```


```
set(gca,'XTickLabelRotation',90);
set(gca,'YTick',1:44);
set(gca,'YTickLabel',labels);
set(gca,'YTickLabelRotation',0);
```

```
labels =
```

```
[]
```

```
nex =
```

```
44
```




First, it is really easy to see something that was evident: the diagonal is zero as it represents the distance between the same utterances. But also, we can observe that the distances between repetitions of the same speaker are really low. We can also observe that the upper-left square corresponds to the lower distances meaning that in general, between the male speakers there are more similarities than between female speakers.

Linkage clustering

Now the objective is to achieve a hierarchical clustering over the Global distance matrix. To do this we use the Linkage Matlab function over the distance matrix and compute the dendrogram with the outputted data.

```
Clusters = linkage(D,'complete');
```

```
[H,T,outperm] = dendrogram(Clusters, 0);
labels = labels(outperm,:);
set(gca,'XTick',1:44);
set(gca,'XTickLabel',labels);
set(gca,'XTickLabelRotation',90);
```



We can observe that the classification of the digits spoken by the same speaker is generally well done, as in the first clustering level, same digits from the same speaker are in the same cluster, with some exceptions: numbers 9 and 8. Also, in the second clustering level we can see that some digits appear together for both speakers as for the digits 2 and 3. The last appreciation will be the fact that man and women are separated in big clusters.

Train

```
groups = cell(4);
character = cell(4);
count = 1;
groups_concatenate = [];
for i = 1:size(tidigits,2)
    if tidigits{i}.digit == '7'
        groups{count} = mfcc(tidigits{i}.samples)';
        groups_concatenate =
            vertcat(groups_concatenate,groups{count});
        character{count} = tidigits{i}.digit;
        count = count + 1;
    end
end
```

```

end

D = zeros(size(groups,2),size(groups,2));
for i = 1:size(groups,2)
    for j = 1:size(groups,2)
        a = groups{i};
        b = groups{j};
        local_dist = localDistances(a, b);
        D(i,j) = dtw(local_dist);
    end
end

componentsNr = [4 8 16 32];
% data = MFCCs{1};
% data = tidigits{1}.samples';
data_t = groups_concatenate;
data = MFCCs_concatenate;

size(data)

options = statset('Display','final');
GMModel = fitgmdist(data_t,componentsNr(1),'Options',options)

% h = ezcontour(@(x,y)pdf(obj,[x y]),[-8 6],[-8 6]);
P = posterior(GMModel,groups{1});
size(P)

figure
subplot(2,1,1)
imagesc(groups{1}')
subplot(2,1,2)
imagesc(P')

ans =

    3883    13

18 iterations, log-likelihood = -18698.7


GMModel =

Gaussian mixture distribution with 4 components in 13 dimensions

ans =

    92    4

```



```

componentsNr = [4 8 16 32];
% data = MFCCs{4};
% data = tidigits{1}.samples';
% data = groups_concatenate;
data = MFCCs_concatenate;
data_t = groups_concatenate;
size(data)

options = statset('Display','final');
GMModel = fitgmdist(data,componentsNr(3),'Options',options)

% h = ezcontour(@(x,y)pdf(obj,[x y]),[-8 6],[-8 6]);
groups_7 = cell(4);
character2 = cell(4);
count = 1;
groups_concatenate_7 = [];
for i = 1:size(tidigits,2)
    if tidigits{i}.digit == '7'
        groups_7{count} = mfcc(tidigits{i}.samples)';
        groups_concatenate_7 =
            vertcat(groups_concatenate,groups_7{count});
        character{count} = tidigits{i}.digit;
        count = count + 1;
    end
end
% groups_concatenate_7

```



```
size(groups_concatenate_7)
mfcc_rows = size(groups_concatenate_7, 1) / 4;
mfcc_cols = size(groups_concatenate_7, 2) / 1;
test1 = groups_concatenate_7(1:mfcc_rows, 1:mfcc_cols);
% P = posterior(GMMModel,data_t);
P = posterior(GMMModel,test1);
% P
size(P)
```

```
figure
subplot(2,1,1)
imagesc(test1')
subplot(2,1,2)
imagesc(P')
```

```
ans =
```

```
3883      13
```

```
Warning: Failed to converge in 100 iterations for gmdistribution with
16
components
100 iterations, log-likelihood = -190562
```

```
GMMModel =
```


```
Gaussian mixture distribution with 16 components in 13 dimensions
```

```
ans =
```

```
484      13
```

```
ans =
```

```
121      16
```



Published with MATLAB® R2016b