

Feature Extraction and Comparing Sequences

DT2119 Speech and Speaker Recognition

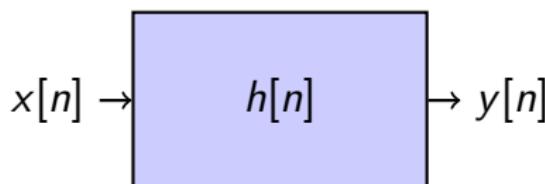
Giampiero Salvi

KTH/CSC/TMH giampi@kth.se

VT 2017

Linear Time-Invariant Systems

Time Domain

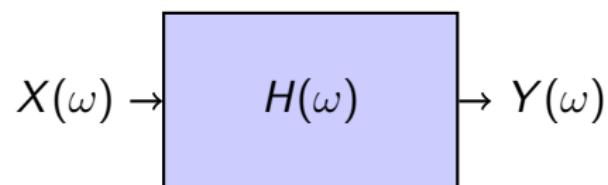


Convolution

$$\begin{aligned}y[n] &= x[n] * h[n] \\&= \sum_{k=-\infty}^{\infty} x[k]h[n-k]\end{aligned}$$

$h[n]$ Impulse Response

Frequency Domain



Multiplication

$$Y(\omega) = X(\omega)H(\omega)$$

$H(\omega)$ Transfer Function

Fourier Transforms

Fourier transform of continuous signals

$$X(\omega) = \int_{-\infty}^{\infty} x(t)e^{j\omega t} dt$$

Fourier transform of discrete signals

$$X(\omega) = \sum_{k=-\infty}^{\infty} x[k]e^{j\omega k}$$

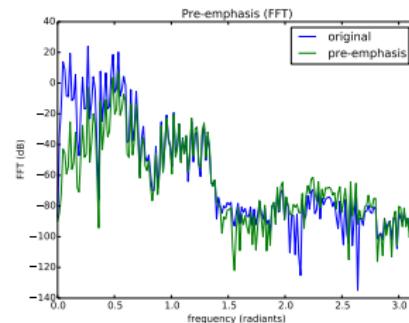
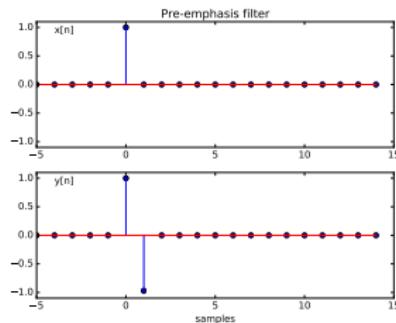
Discrete Fourier Transform (implemented as FFT)

$$X[n] = \sum_{k=-\infty}^{\infty} x[k]e^{j2\pi \frac{n}{N} k}$$

Properties of Fourier Transform

Time domain	Frequency domain	
$x[n] * h[n]$	\Leftrightarrow	$X(\omega)H(\omega)$

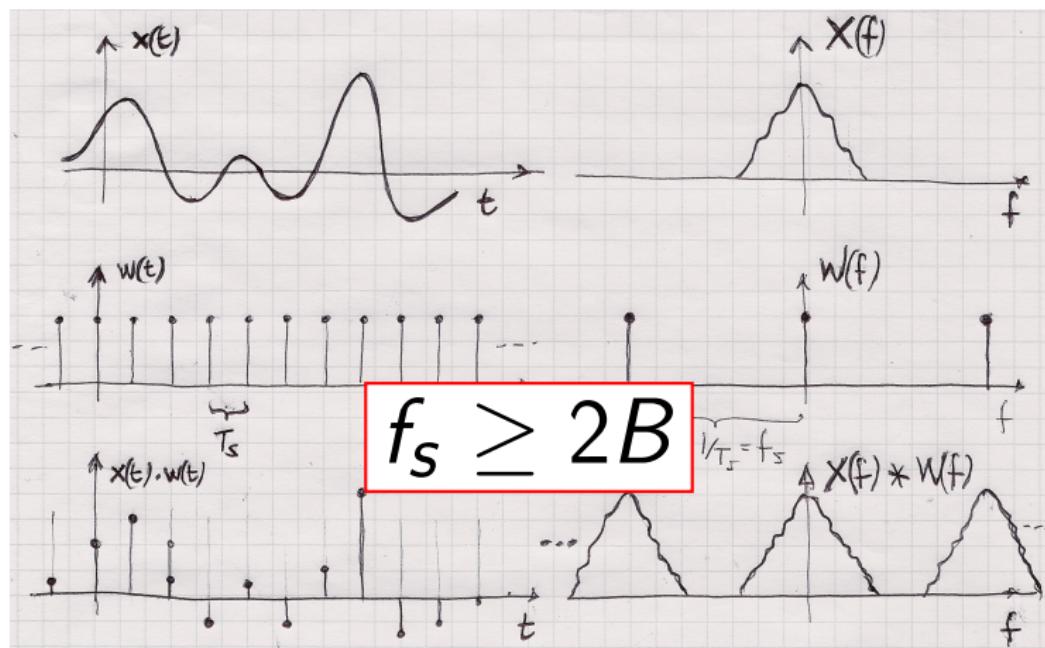
Example: any filter (pre-emphasis, moving average...)



Properties of Fourier Transform

Time domain	\iff	Frequency domain
$x[n]w[n]$		$X(\omega) * W(\omega)$

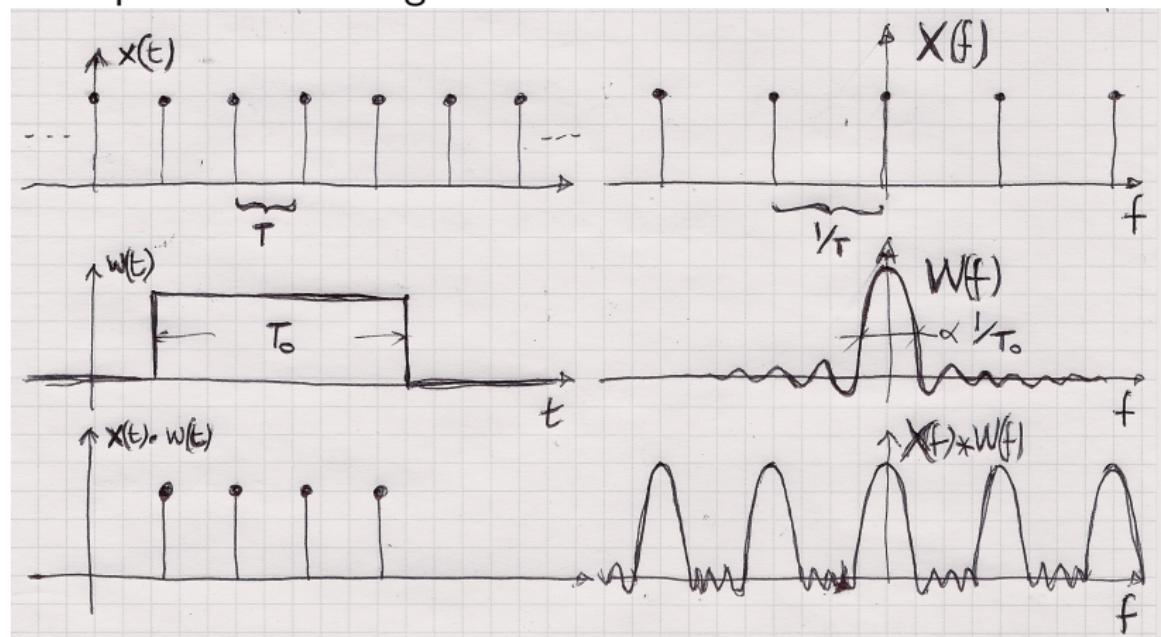
Example 1: sampling



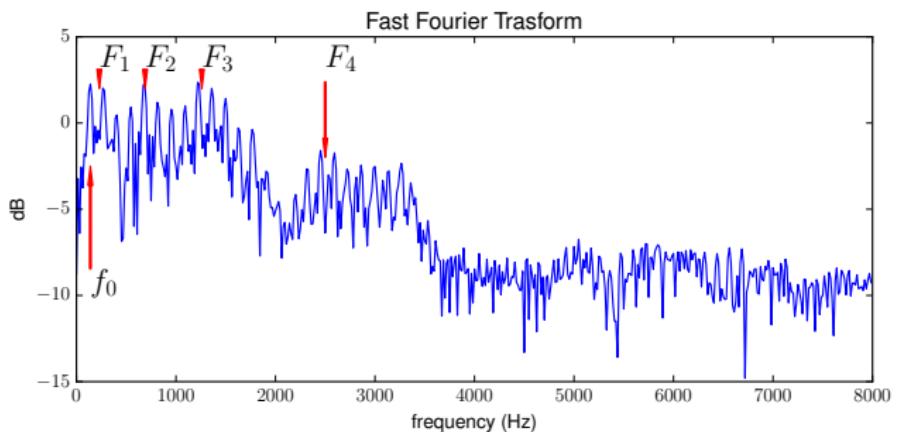
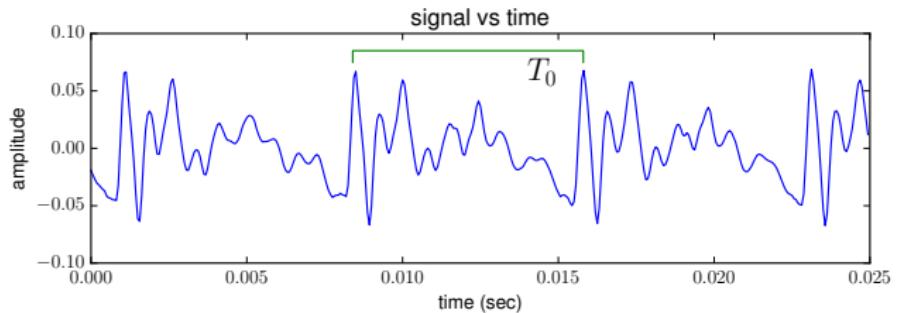
Properties of Fourier Transform

Time domain	Frequency domain
$x[n]w[n]$	$\Leftrightarrow X(\omega) * W(\omega)$

Example 2: windowing

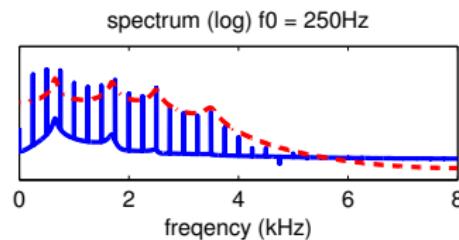
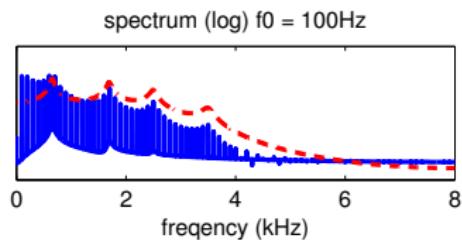


Example: vowel

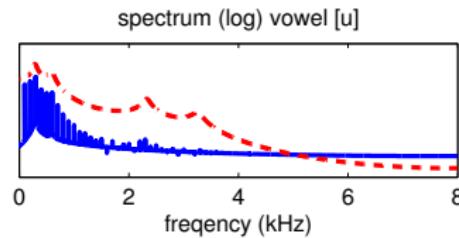
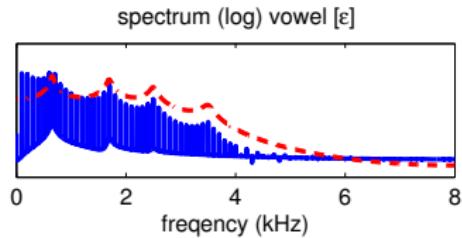


F_0 and Formants

- ▶ Varying F_0 (vocal fold oscillation rate)



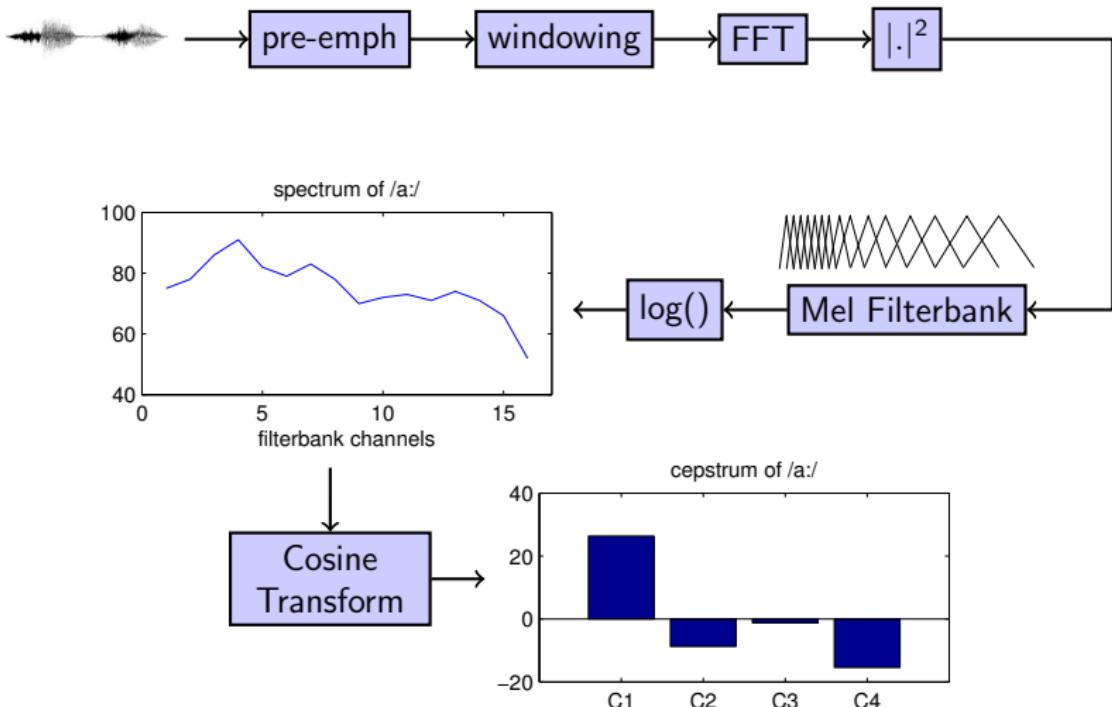
- ▶ Varying Formants (vocal tract shape)



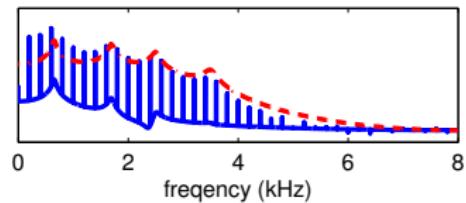
Mel Frequency Cepstrum Coefficients

- ▶ *de facto* standard in ASR (before Deep Learning)
- ▶ imitate aspects of auditory processing
- ▶ does not assume all-pole model of the spectrum
- ▶ uncorrelated: easier to model statistically

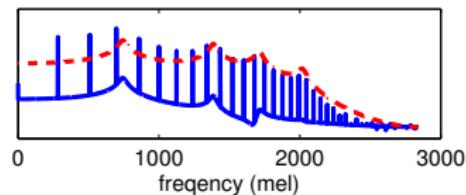
MFCCs Calculation



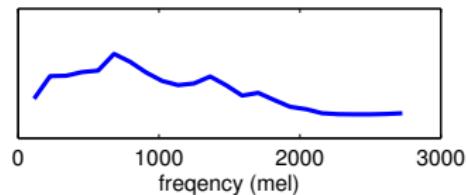
Mel Frequency Cepstral Coefficients



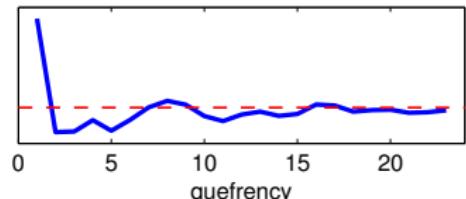
Linear to Mel frequency



Filterbank ($\sim 20\text{-}25$ filters) + $\log()$

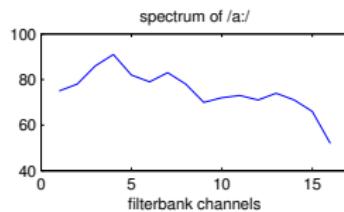


Discrete Cosine Transform

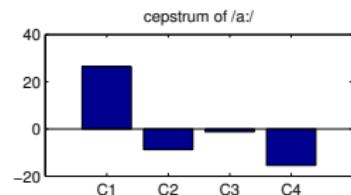
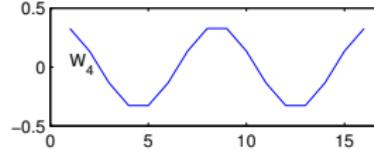
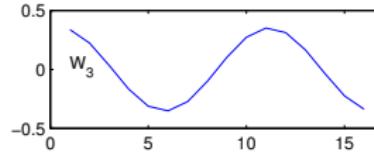
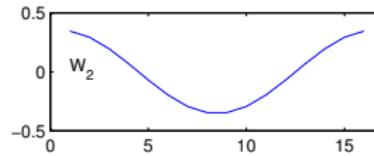
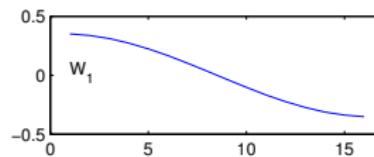
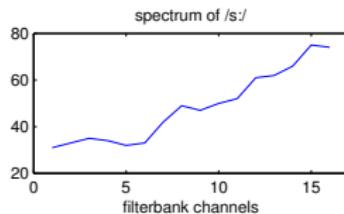


MFCC: Cosine Transform

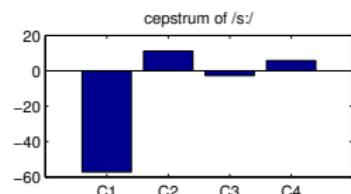
$$C_j = \sqrt{\frac{2}{N}} \sum_{i=1}^N A_i \cos\left(\frac{j\pi(i-0.5)}{N}\right)$$



A_i



C_j



MFCC Rationale

- ▶ signals combined in a convolutive way: $a[n] * b[n] * c[n]$
- ▶ in the spectral domain: $A(z)B(z)C(z)$
- ▶ taking the log: $\log(A(z)) + \log(B(z)) + \log(C(z))$
- ▶ to analyse the different contribution perform Fourier transform (DCT if not interested in phase information).
- ▶ Terminology:
 - ▶ frequency vs quefrency
 - ▶ spectrum vs cepstrum
 - ▶ filter vs lifter
 - ▶ ...

MFCC Advantages [1]

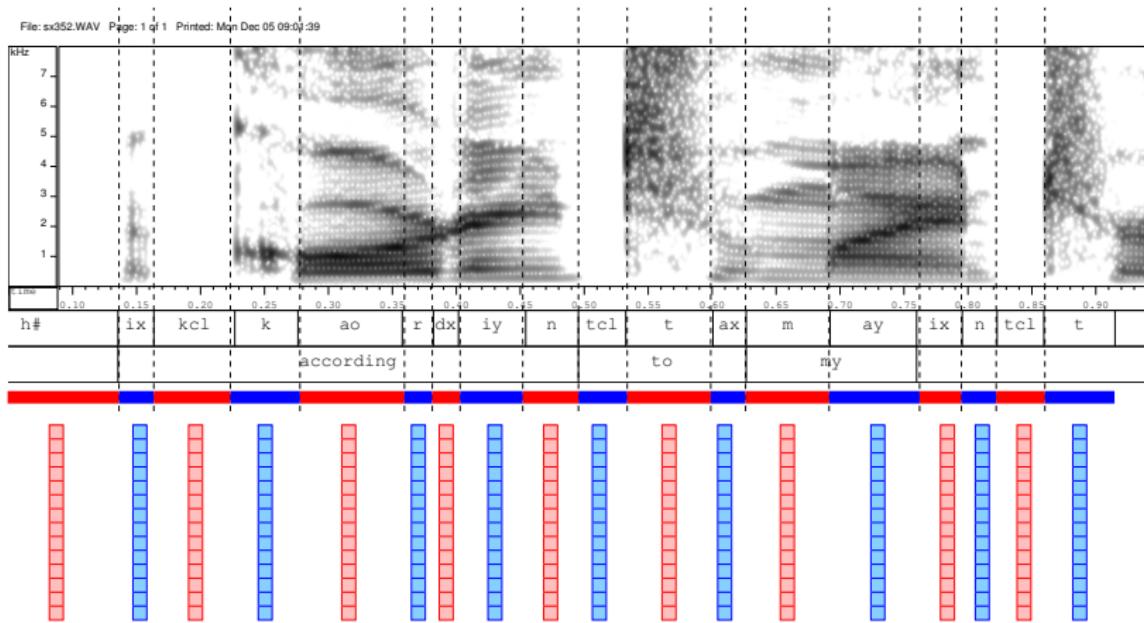
- ▶ fairly uncorrelated coefficients (simpler statistical models)
- ▶ high phonetic discrimination (empirically shown)
- ▶ do not assume all-pole model
- ▶ low number of coeff. enough to capture coarse structure of spectrum
- ▶ Cepstral Mean Subtraction corresponds to channel removal

[1] B. Bogert, M. Healy, and J. Tukey. "The Quefrency Alanysis of Time Series for Echoes: Cepstrum, Pseudo-autocovariance, Cross-Cepstrum and Saphe Cracking". In: *Proc. Symp. Time Series Analysis*. Ed. by M. Rosemblatt. John Wiley & Sons, 1963, pp. 209–243

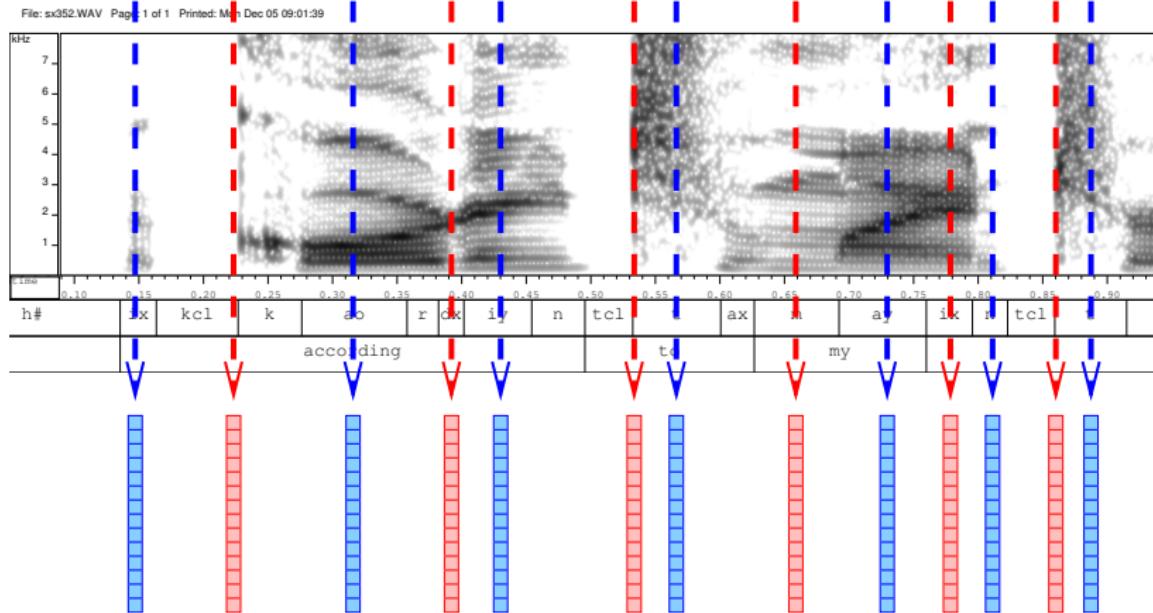
MFCCs: typical values

- ▶ 12 Coefficients C1–C12
- ▶ Energy (could be C0)
- ▶ Delta coefficients (derivatives in time)
- ▶ Delta-delta (second order derivatives)
- ▶ total: 39 coefficients per frame (analysis window)

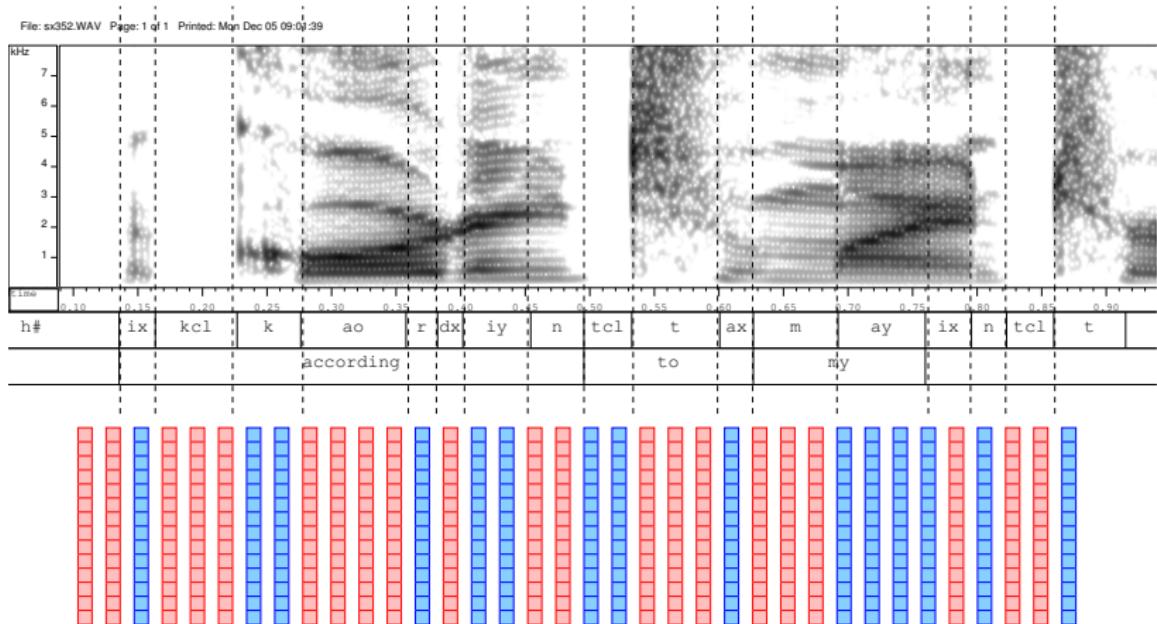
Segment-Based Processing



Landmark-Based Processing



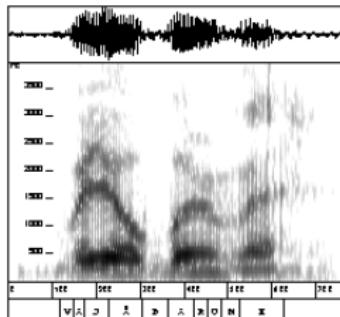
Frame-Based Processing



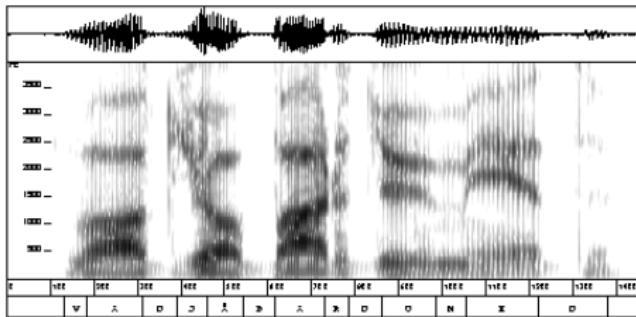
Frame-wise distance metrics

distance	$d(x, y)$
city block:	$\sum_i x_i - y_i $
Euclidean:	$\sqrt{\sum_i (x_i - y_i)^2}$
Mahalanobis:	$\sum_i (x_i - \mu_y)^2 / \sigma_y$
probability function:	$f(X = x \mu_y, \Sigma_y)$
artificial neural networks:	$f(\sum_i w_i x_i - \theta)$

Comparing Utterances



Va jobbaru me



Vad jobbar du med

“What is your occupation”
 (“What work you with”)

Combining frame-wise scores into utterance scores

Template Matching

- ▶ oldest technique
- ▶ simple comparison of template patterns
- ▶ compensate for varying speech rate (Dynamic Programming)

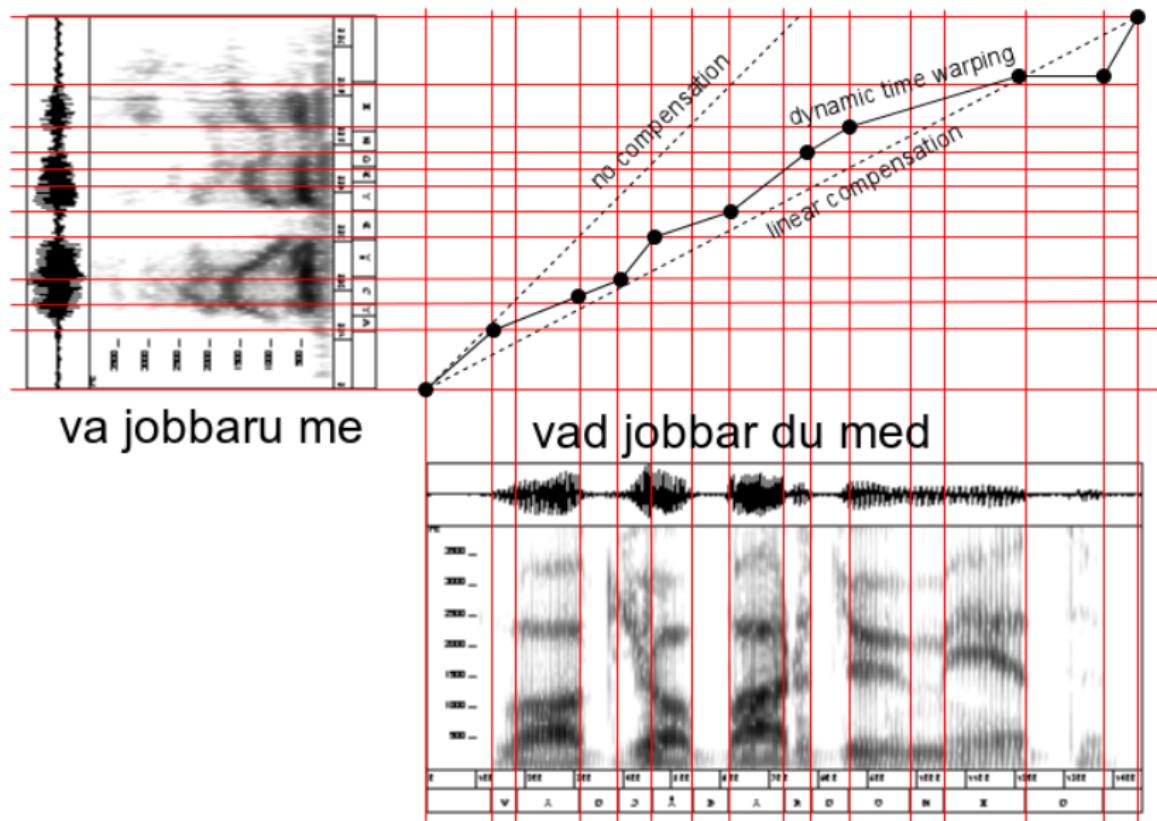
Hidden Markov Models (HMMs)

- ▶ most used technique
- ▶ models of segmental structure of speech
- ▶ recognition by Viterbi search (Dynamic Programming)

Template Matching: Why?

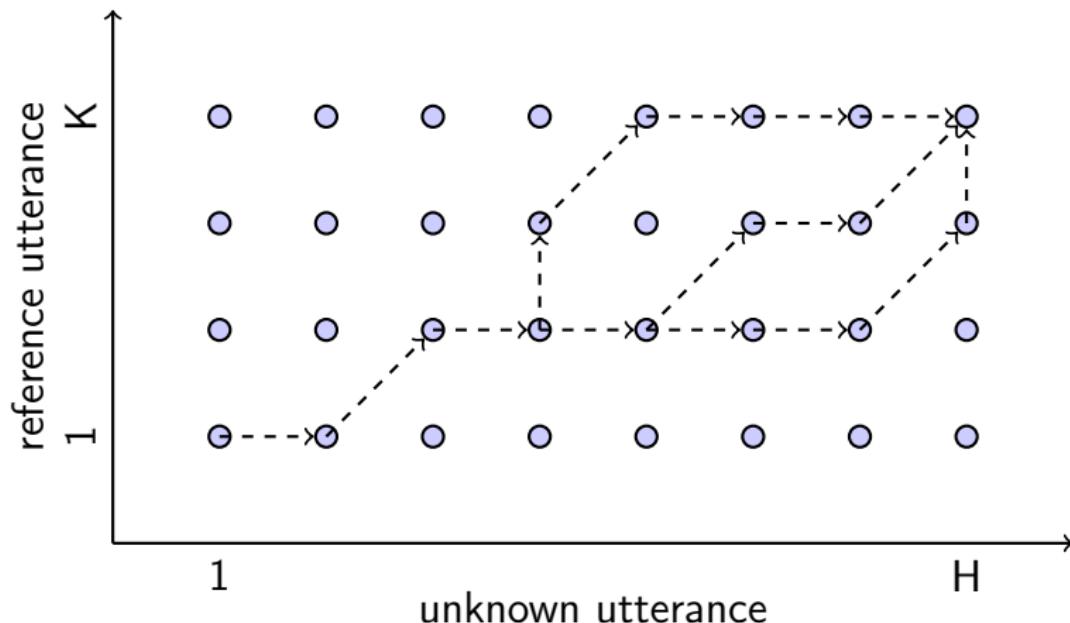
1. Historical: first technique used in ASR
2. Pedagogical: explain the problem and Dynamic Programming
3. Also called Dynamic Time Warping (DTW)

Template Matching



Dynamic Programming

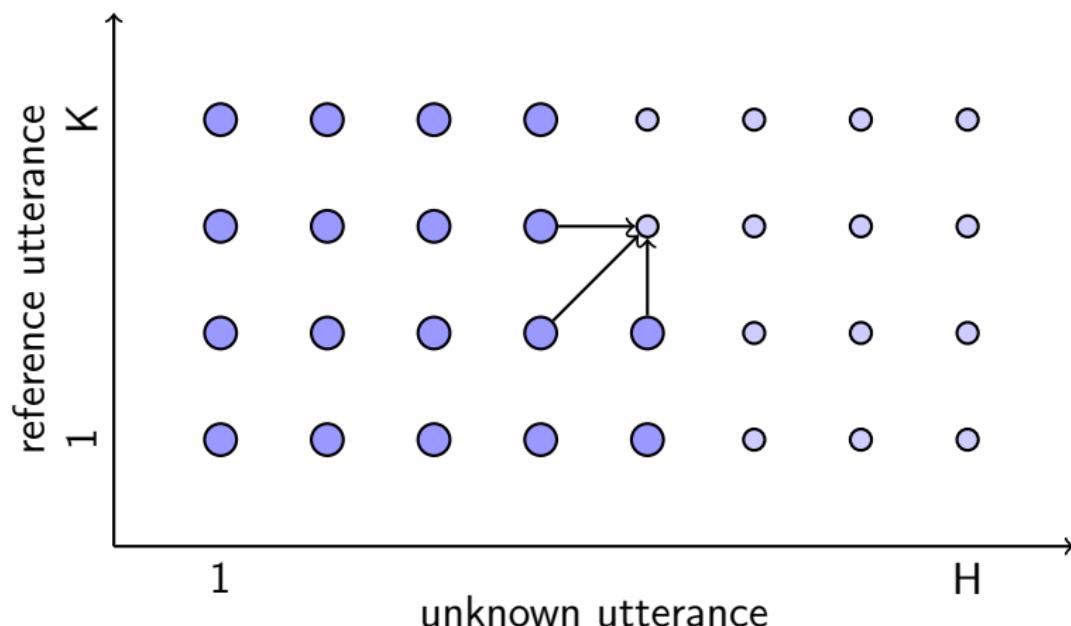
- ▶ compare any possible alignment
 - ▶ problem: exponential with H and K!



Dynamic Programming

Dynamic Time Warping (DTW) algorithm

```
1: for  $h = 1$  to  $H$  do
2:   for  $k = 1$  to  $K$  do
3:      $AccD[h, k] = LocD[h, k] + \min(AccD[h - 1, k], AccD[h - 1, k - 1], AccD[h, k - 1])$ 
```



DP Example: Spelling

- ▶ observations are letters
- ▶ local distance: 0 (same letter), 1 (different letter)
- ▶ Unknown utterance: ALLDRIG
- ▶ Reference1: ALDRIG
- ▶ Reference2: ALLTID
- ▶ Problem: find closest match

Distance char-by-char:

- ▶ ALLDRIG–ALDRIG = 5
- ▶ ALLDRIG–ALLTID = 4

DP Example: Solution

$LocD[h,k] =$

G	1	1	1	1	1	1	1	0
I	1	1	1	1	1	0	1	
R	1	1	1	1	0	1	1	
D	1	1	1	0	1	1	1	
L	1	0	0	1	1	1	1	
A	0	1	1	1	1	1	1	
	A	L	L	D	R	I	G	

$AccD[h,k] =$

G	5	4	4	3	2	1	0
I	4	3	3	2	1	0	1
R	3	2	2	1	0	1	2
D	2	1	1	0	1	2	3
L	1	0	0	1	2	3	4
A	0	1	2	3	4	5	6
	A	L	L	D	R	I	G

Distance ALLDRIG–ALDRIG: $AccD[H,K] = 0$

Distance ALLDRIG–ALLTID? (5min)

DP Example: Solution

$LocD[h,k] =$

D	1	1	1	0	1	1	1	1
I	1	1	1	1	1	0	1	
T	1	1	1	1	1	1	1	
L	1	0	0	1	1	1	1	
L	1	0	0	1	1	1	1	
A	0	1	1	1	1	1	1	
	A	L	L	D	R	I	G	

$AccD[h,k] =$

D	5	3	3	2	3	3	3	
I	4	2	2	2	2	2	3	
T	3	1	1	1	2	3	4	
L	2	0	0	1	2	3	4	
L	1	0	0	1	2	3	4	
A	0	1	2	3	4	5	6	
	A	L	L	D	R	I	G	

Distance ALLDRIG–ALDRIG: $AccD[H,K] = 0$

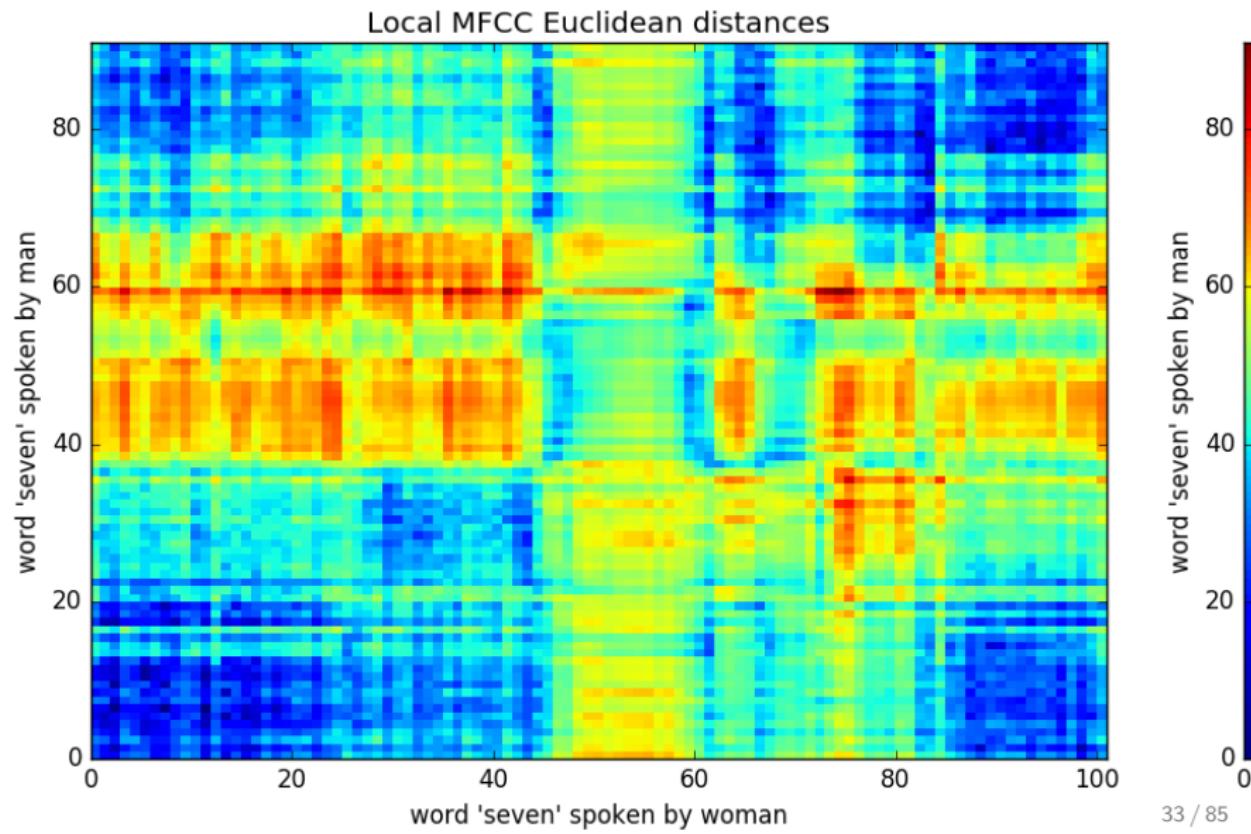
Distance ALLDRIG–ALLTID: $AccD[H,K] = 3$

Best path: Backtracking

Sometimes we want to know the path

1. at each point $[h,k]$ remember the minimum distance predecessor (back pointer)
2. at the end point $[H,K]$ follow the back pointers until the start

Real Example from Lab1



Properties of Template Matching

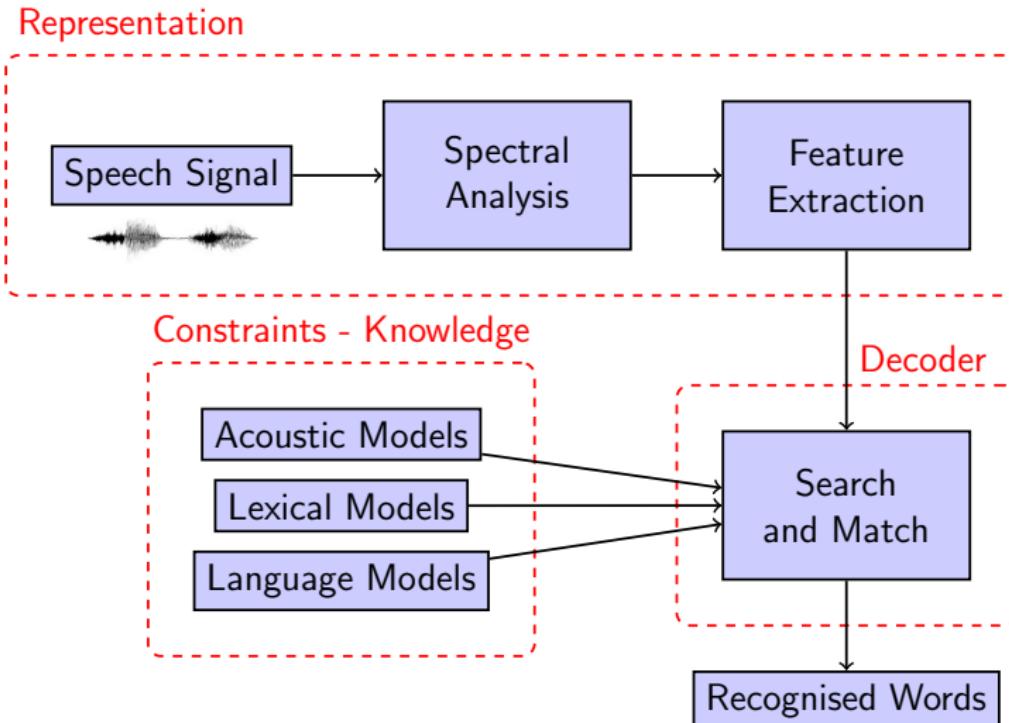
Pros:

- + No need for phonetic transcriptions
- + within-word co-articulation for free
- + high time resolution

Cons:

- cross-word co-articulation not modelled
- needs word segmentation (isolated words)
- requires recordings of every word
- not easy to model variation
- does not scale up with vocabulary size

Components of ASR System



Assuming you know:

- ▶ Probability theory
- ▶ Basic probability distributions
- ▶ Basic machine learning

Bayes' Rule

if

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

then

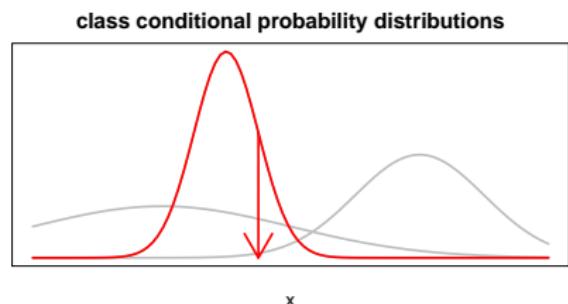
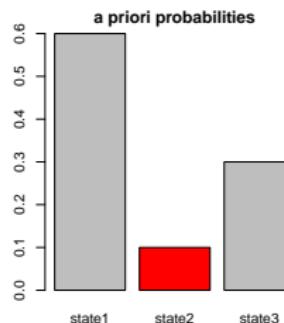
$$P(A \cap B) = P(A|B)P(B) = P(B|A)P(A)$$

and

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

The Probabilistic Model of Classification

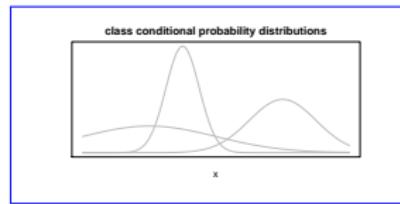
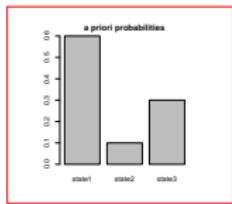
- ▶ “Nature” assumes one of c states ω_j with *a priori* probability $P(\omega_j)$
- ▶ When in state ω_j , “nature” emits observations $\hat{\mathbf{x}}$ with distribution $p(\mathbf{x}|\omega_j)$



Problem

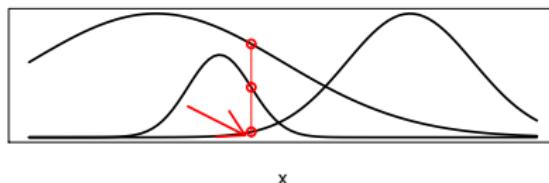
- ▶ If I observe $\hat{\mathbf{x}}$ and I know $P(\omega_j)$ and $p(\mathbf{x}|\omega_j)$ for each j
- ▶ what can I say about the state of “nature” ω_j ?

Bayes decision theory

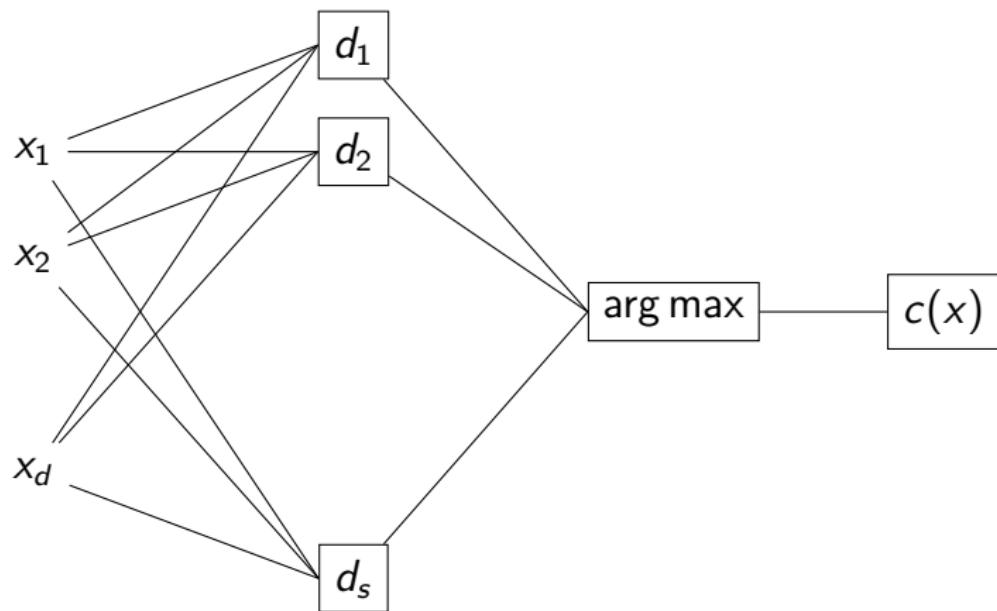


$$P(\omega_j | \mathbf{x}) = \frac{p(\mathbf{x} | \omega_j) P(\omega_j)}{p(\mathbf{x})}$$

posterior probabilities



Classifiers: Discriminant Functions



$$d_i(\mathbf{x}) = p(\mathbf{x}|\omega_i) P(\omega_i)$$

Classifiers: Decision Boundaries

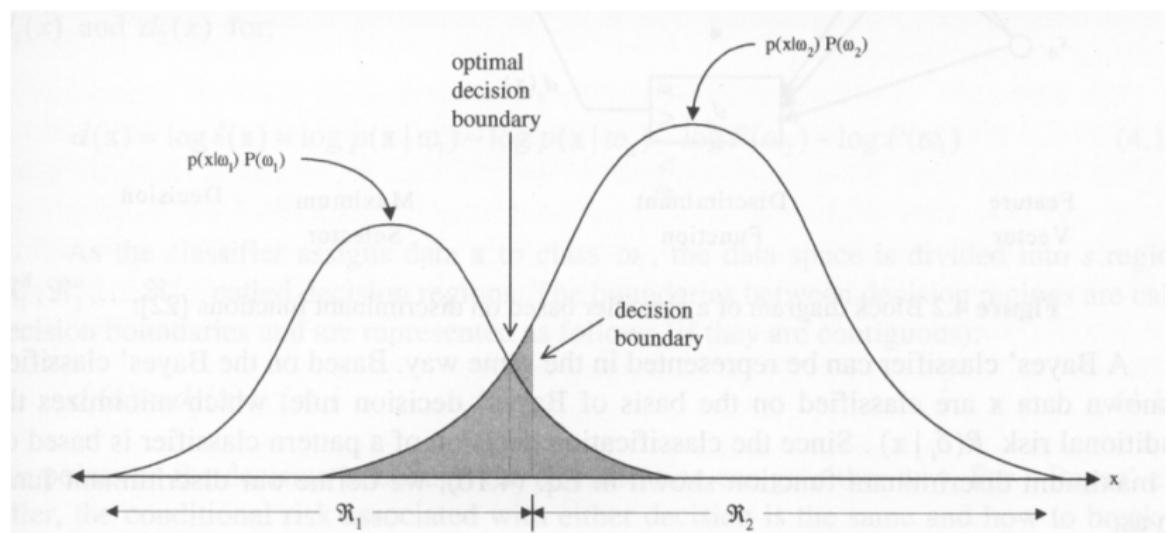


Figure from Huang, Acero, Hon.

Decision Boundaries in Two Dimensions

<https://github.com/giampierosalvi/GaussianDecisionBoundaries>

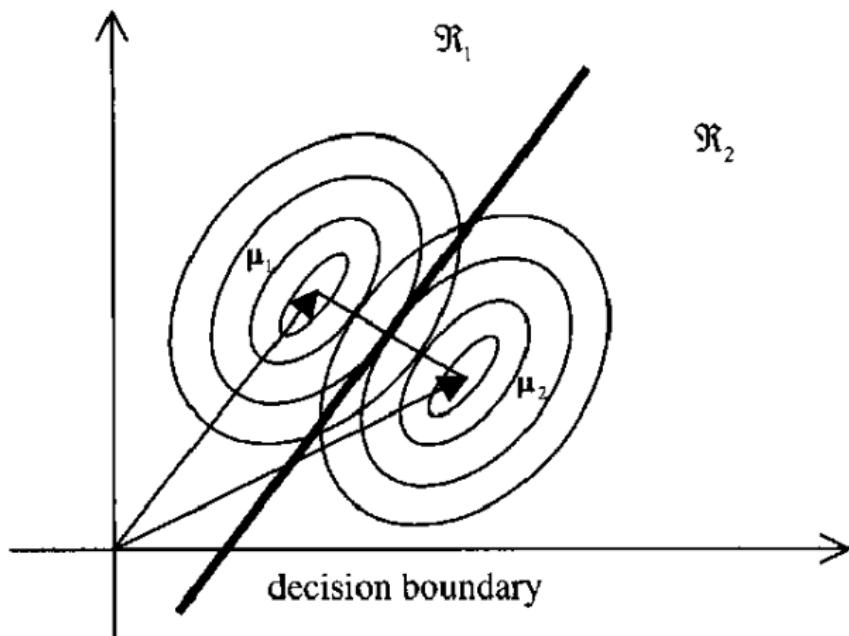


Figure from Huang, Acero, Hon.

Bayes' Rule and Speech Recognition

$A = \text{words}$, $B = \text{sounds}$:

- ▶ During training we know the words and can compute $P(\text{sounds}|\text{words})$ using frequentist approach (repeated observations)
- ▶ during recognition we want
 $\hat{\text{words}} = \arg \max P(\text{words}|\text{sounds})$
- ▶ using Bayes' rule:

$$P(\text{words}|\text{sounds}) = \frac{P(\text{sounds}|\text{words})P(\text{words})}{P(\text{sounds})}$$

where

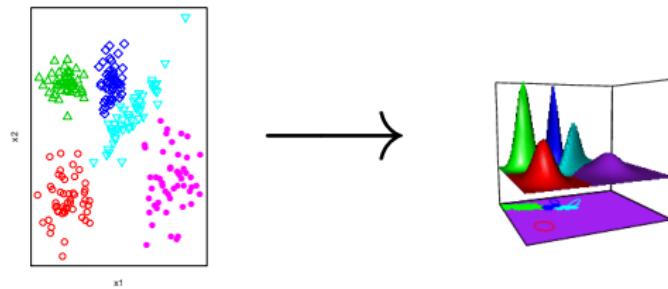
$P(\text{words})$: *a priori* probability of the words (Language Model)

$P(\text{sounds})$: *a priori* probability of the sounds (constant, can be ignored)

Estimation Theory

- ▶ so far we assumed we know $P(\omega_j)$ and $p(\mathbf{x}|\omega_j)$
- ▶ how can we obtain them from collections of data?
- ▶ this is the subject of Estimation Theory

Parameter estimation

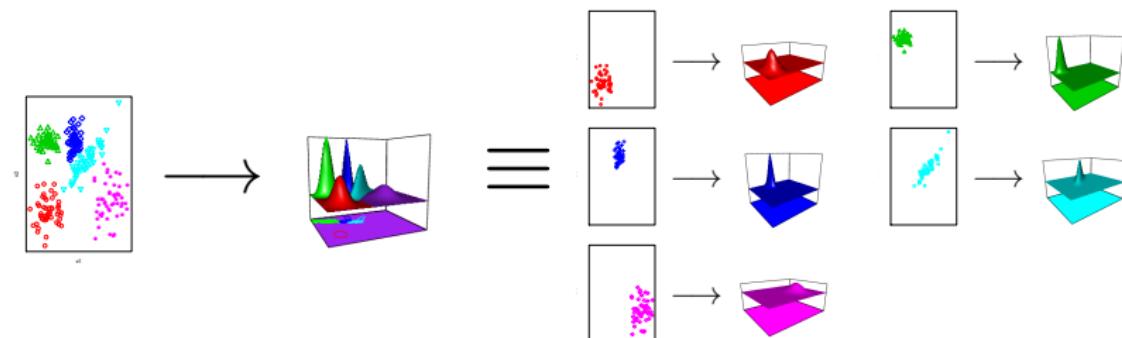


Assumptions:

- ▶ samples from class ω_i do not influence estimate for class ω_j , $i \neq j$
- ▶ samples from the same class are independent and identically distributed (i.i.d.)

Parameter estimation (cont.)

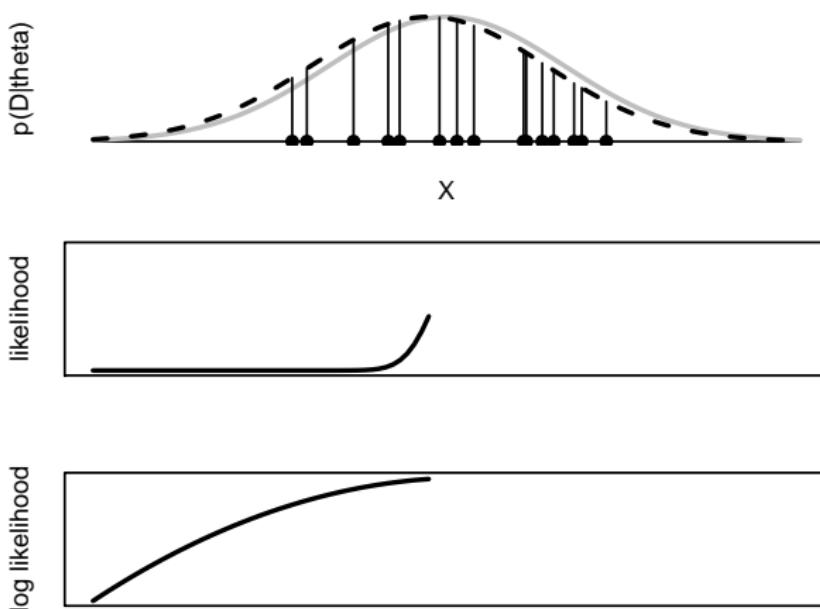
- ▶ class independence assumption:



- ▶ Maximum likelihood estimation
- ▶ Maximum a posteriori estimation
- ▶ Bayesian estimation

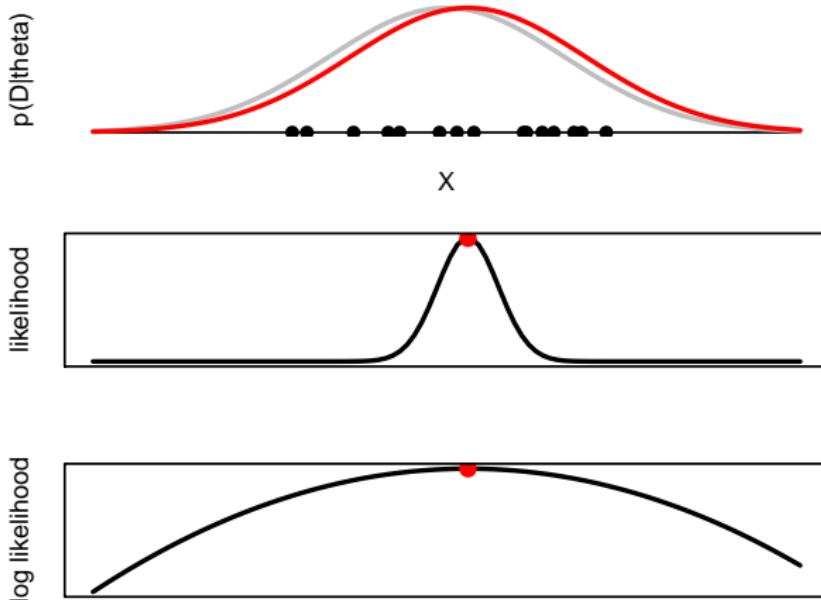
Maximum likelihood estimation

- ▶ Find parameter vector $\hat{\theta}$ that maximises $p(\mathcal{D}|\theta)$ with $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- ▶ i.i.d. $\rightarrow p(\mathcal{D}|\theta) = \prod_{k=1}^n p(\mathbf{x}_k|\theta)$



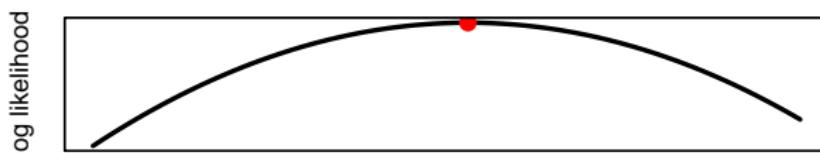
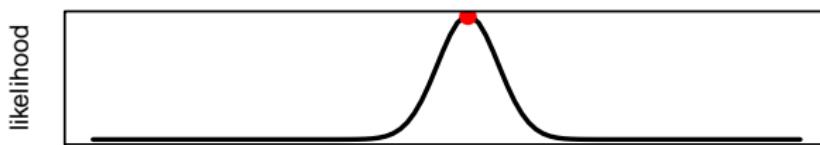
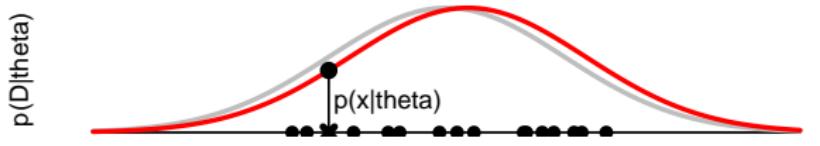
Maximum likelihood estimation

- ▶ Find parameter vector $\hat{\theta}$ that maximises $p(\mathcal{D}|\theta)$ with $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- ▶ i.i.d. $\rightarrow p(\mathcal{D}|\theta) = \prod_{k=1}^n p(\mathbf{x}_k|\theta)$



Maximum likelihood estimation

- ▶ Find parameter vector $\hat{\theta}$ that maximises $p(\mathcal{D}|\theta)$ with $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- ▶ i.i.d. $\rightarrow p(\mathcal{D}|\theta) = \prod_{k=1}^n p(\mathbf{x}_k|\theta)$



ML estimation of Gaussian mean

$$N(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left[-\frac{(x-\mu)^2}{2\sigma^2}\right], \text{ with } \theta = \{\mu, \sigma^2\}$$

Log-likelihood of data (i.i.d. samples):

$$\log P(\mathcal{D}|\theta) = \sum_{i=1}^N \log N(x_i|\mu, \sigma^2) = -N \log (\sqrt{2\pi}\sigma) - \sum_{i=1}^N \frac{(x_i - \mu)^2}{2\sigma^2}$$

$$0 = \frac{d \log P(\mathcal{D}|\theta)}{d\mu} = \sum_{i=1}^N \frac{(x_i - \mu)}{\sigma^2} = \frac{\sum_{i=1}^N x_i - N\mu}{\sigma^2} \iff$$

$$\hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i$$

ML estimation of Gaussian parameters

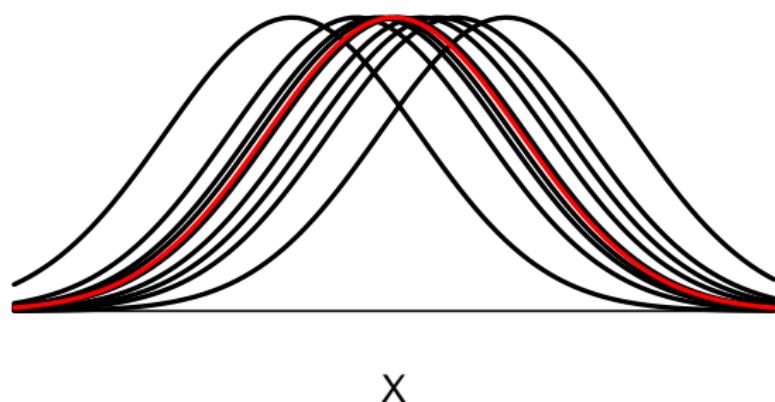
$$\hat{\mu} = \frac{1}{N} \sum_{i=1}^N x_i$$

$$\hat{\sigma}^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \hat{\mu})^2$$

- ▶ same result by minimizing the sum of square errors!
- ▶ but we make assumptions explicit

Problem: few data points

10 repetitions with 5 points each



Maximum a Posteriori Estimation

$$\hat{\mu}, \hat{\sigma}^2 = \arg \max_{\mu, \sigma^2} \left[\prod_{i=1}^N P(x_i | \mu, \sigma^2) P(\mu, \sigma^2) \right]$$

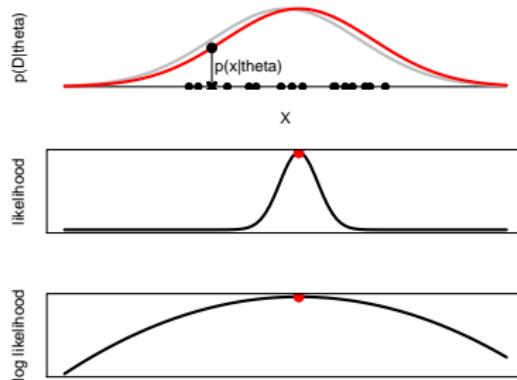
where the prior $P(\mu, \sigma^2)$ needs a nice mathematical form for closed solution

$$\begin{aligned}\hat{\mu}_{\text{MAP}} &= \frac{N}{N + \gamma} \hat{\mu}_{\text{ML}} + \frac{\gamma}{N + \gamma} \delta \\ \hat{\sigma}_{\text{MAP}}^2 &= \frac{N}{N + 3 + 2\alpha} \hat{\sigma}_{\text{ML}}^2 + \frac{2\beta + \gamma(\delta + \hat{\mu}_{\text{MAP}})^2}{N + 3 + 2\alpha}\end{aligned}$$

where $\alpha, \beta, \gamma, \delta$ are parameters of the prior distribution

ML, MAP and Point Estimates

- Both ML and MAP produce point estimates of θ
- Assumption: there is a **true** value for θ
- advantage: once $\hat{\theta}$ is found, everything is known



Overfitting

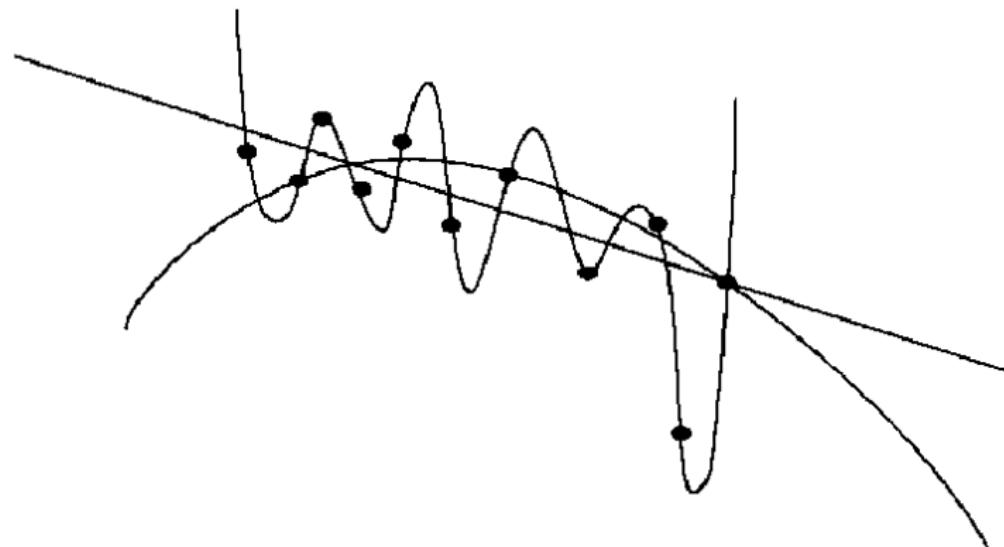


Figure from Huang, Acero, Hon.

Overfitting: Phoneme Discrimination

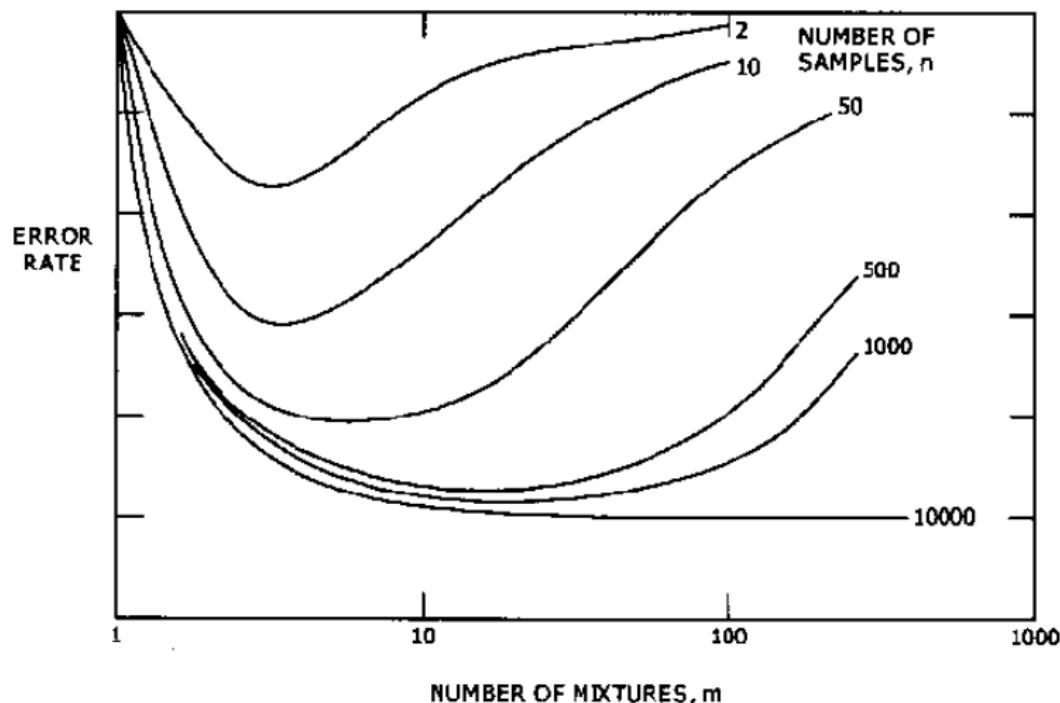


Figure from Huang, Acero, Hon.

Bayesian estimation

- ▶ Consider θ as a random variable
- ▶ characterize θ with the posterior distribution $P(\theta|\mathcal{D})$ given the data

$$\text{ML: } \mathcal{D} \rightarrow \hat{\theta}_{\text{ML}}$$

$$\text{MAP: } \mathcal{D}, P(\theta) \rightarrow \hat{\theta}_{\text{MAP}}$$

$$\text{Bayes: } \mathcal{D}, P(\theta) \rightarrow P(\theta|\mathcal{D})$$

- ▶ for new data points, instead of $P(\mathbf{x}_{\text{new}}|\hat{\theta}_{\text{ML}})$ or $P(\mathbf{x}_{\text{new}}|\hat{\theta}_{\text{MAP}})$, compute:

$$P(\mathbf{x}_{\text{new}}|\mathcal{D}) = \int_{\theta \in \Theta} P(\mathbf{x}_{\text{new}}|\theta) P(\theta|\mathcal{D}) d\theta$$

Bayesian estimation (cont.)

Pros:

- ▶ better use of the data
- ▶ makes a priori assumptions explicit
- ▶ easily implemented recursively
 - ▶ use posterior $p(\theta|\mathcal{D})$ as new prior
- ▶ reduce overfitting

Cons:

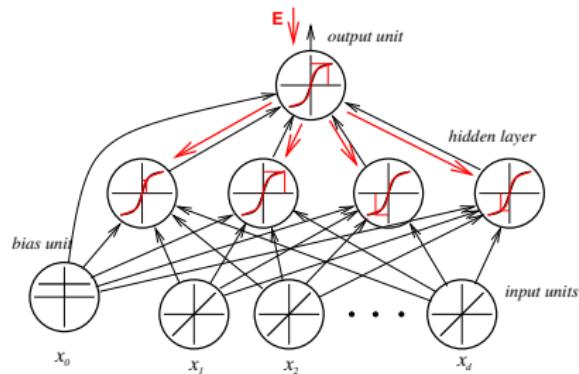
- ▶ definition of noninformative priors can be tricky
- ▶ often requires numerical integration

Other Training Strategies: Discriminative Training

- ▶ Maximum Mutual Information Estimation
- ▶ Minimum Error Rate Estimation
- ▶ Neural Networks

Multi layer neural networks

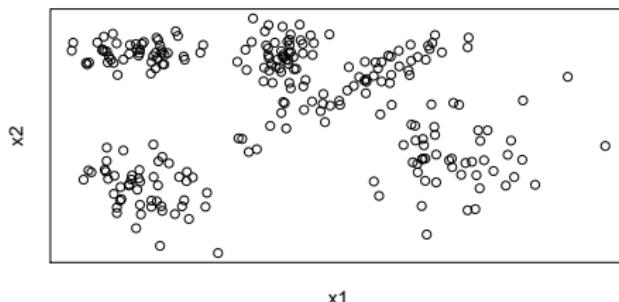
Multi layer
neural networks



- ▶ Backpropagation algorithm

Unsupervised Learning

- ▶ so far we assumed we knew the class ω_i for each data point
- ▶ what if we don't?
- ▶ class independence assumption loses meaning



Vector Quantisation, K-Means

- ▶ describes each class with a centroid
- ▶ a point belongs to a class if the corresponding centroid is closest (Euclidean distance)
- ▶ iterative procedure
- ▶ guaranteed to converge
- ▶ not guaranteed to find the optimal solution
- ▶ used in vector quantization

K-means: algorithm

Data: k (number of desired clusters), n data points \mathbf{x}_i

Result: k clusters

initialization: assign initial value to k centroids \mathbf{c}_i ;

repeat

 assign each point \mathbf{x}_i to closest centroid \mathbf{c}_j ;

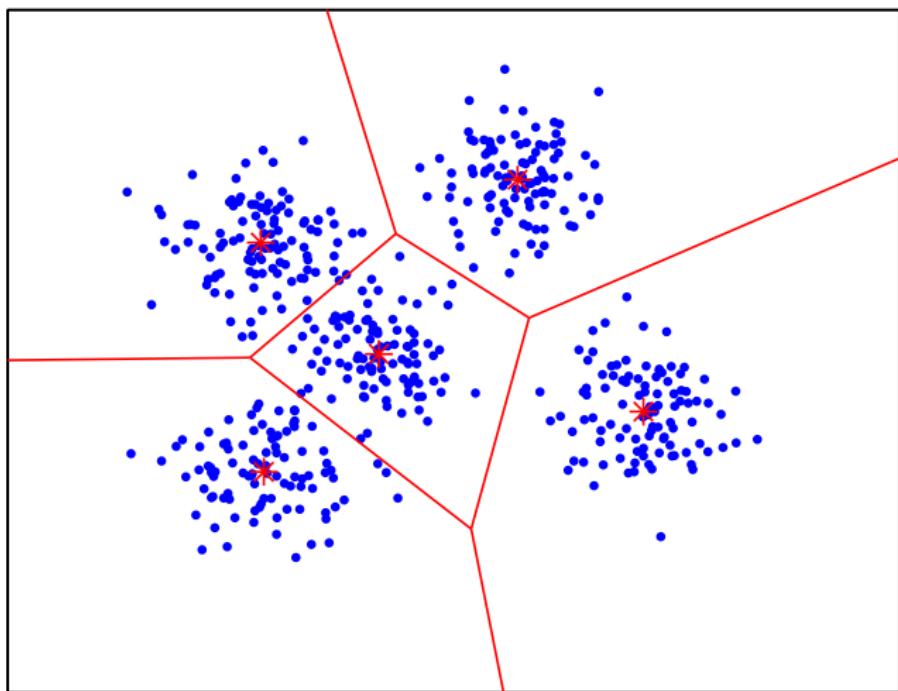
 compute new centroids as mean of each group of points;

until centroids do not change;

return k clusters;

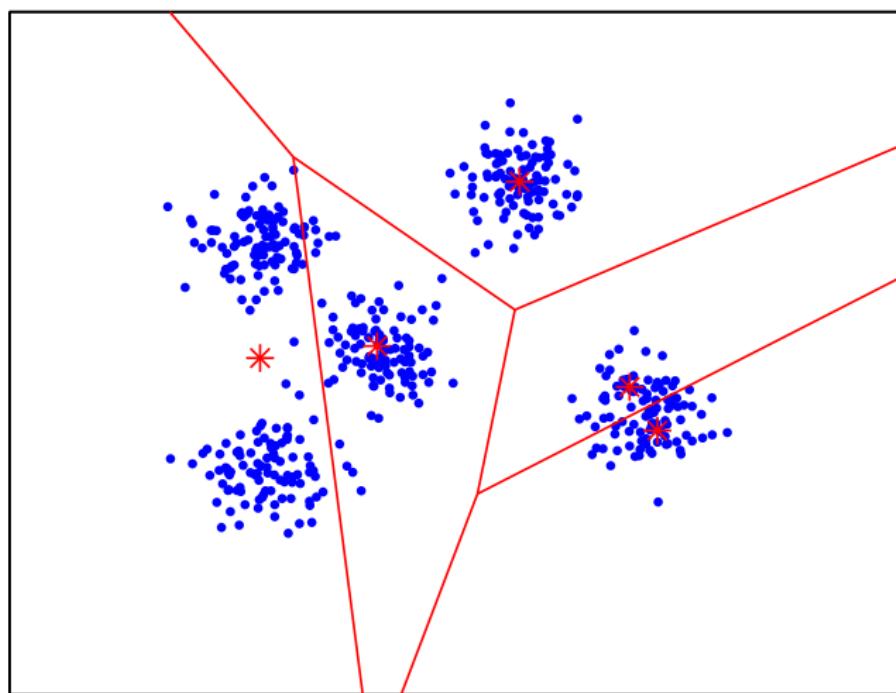
K-means: example

iteration 20, update clusters



K-means: sensitivity to initial conditions

iteration 20, update clusters



Solution: LBG Algorithm

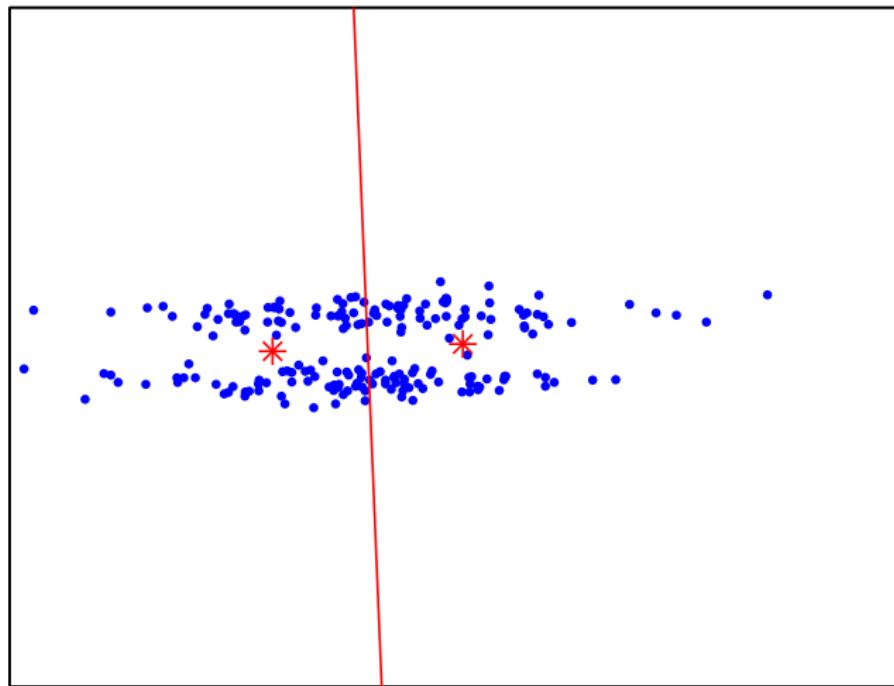
- ▶ Linde–Buzo–Gray
- ▶ start with one centroid
- ▶ adjust to mean
- ▶ split centroid (with ϵ)
- ▶ K-means
- ▶ split again...

K-means: limits of Euclidean distance

- ▶ the Euclidean distance is isotropic (same in all directions in \mathbb{R}^p)
- ▶ this favours spherical clusters
- ▶ the size of the clusters is controlled by their distance

K-means: non-spherical classes

two non-spherical classes



Probabilistic Clustering

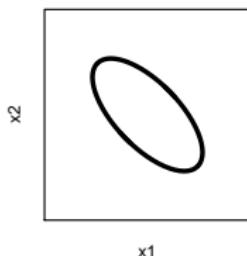
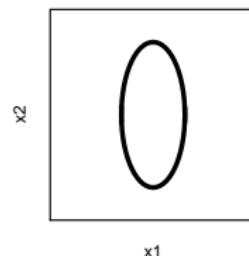
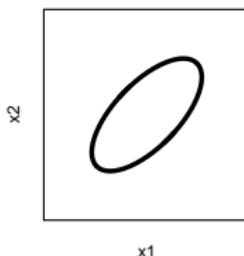
- ▶ model data as a mixture of probability distributions (Gaussian)
- ▶ each distribution corresponds to a cluster
- ▶ clustering corresponds to parameter estimation

Gaussian distributions

$$f_k(\mathbf{x}_i | \mu_k, \Sigma_k) = \frac{\exp\left\{-\frac{1}{2}(\mathbf{x}_i - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_i - \mu_k)\right\}}{(2\pi)^{\frac{p}{2}} |\Sigma_k|^{\frac{1}{2}}}$$

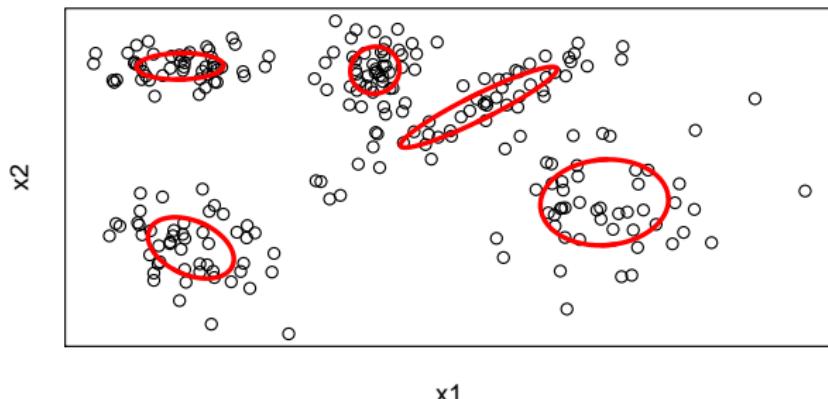
Eigenvalue decomposition of the covariance matrix:

$$\Sigma_k = \lambda_k D_k A_k D_k^T$$



Mixture of Gaussian distributions

Σ_k	Distribution	Volume	Shape	Orientation
λI	Spherical	Equal	Equal	N/A
$\lambda_k I$	Spherical	Variable	Equal	N/A
λDAD^T	Ellipsoidal	Equal	Equal	Equal
$\lambda D_k A D_k^T$	Ellipsoidal	Equal	Equal	Variable
$\lambda_k D_k A D_k^T$	Ellipsoidal	Variable	Equal	Variable
$\lambda_k D_k A_k D_k^T$	Ellipsoidal	Variable	Variable	Variable



Fitting the model

- ▶ given the data $D = \{\mathbf{x}_i\}$
- ▶ given a certain model \mathcal{M} and its parameters θ
- ▶ maximize the model fit to the data as expressed by the likelihood

$$\mathcal{L} = p(D|\theta)$$

Unsupervised Case

- ▶ release class independence assumption:
- ▶ learn the mixture at once
- ▶ problem of missing data
- ▶ solution: Expectation Maximization

Expectation Maximization

Fitting model parameters with missing (**latent**) variables

$$P(\mathbf{x}|\theta) = \sum_{k=1}^K \pi_k P(x|\theta_k),$$

$$\text{with } \theta = \{\pi_1, \dots, \pi_K, \theta_1, \dots, \theta_K\}$$

- ▶ very general idea (applies to many different probabilistic models)
- ▶ augment the data with the missing variables: h_{ik} probability of assignment of each data point x_i to each component of the mixture k
- ▶ optimize the Likelihood of the complete data:

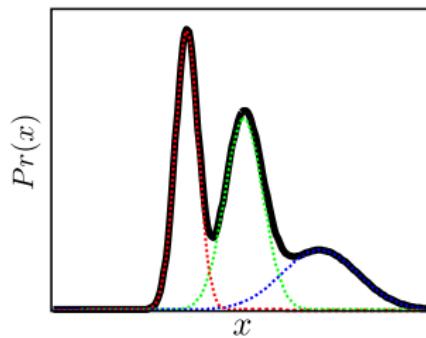
$$P(\mathbf{x}, \mathbf{h}|\theta)$$

Mixture of Gaussians

This distribution is a weight sum of K Gaussian distributions

$$P(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x; \mu_k, \sigma_k^2)$$

where $\pi_1 + \dots + \pi_K = 1$
and $\pi_k > 0$ ($k = 1, \dots, K$).



This model can describe **complex multi-modal** probability distributions by combining simpler distributions.

Mixture of Gaussians

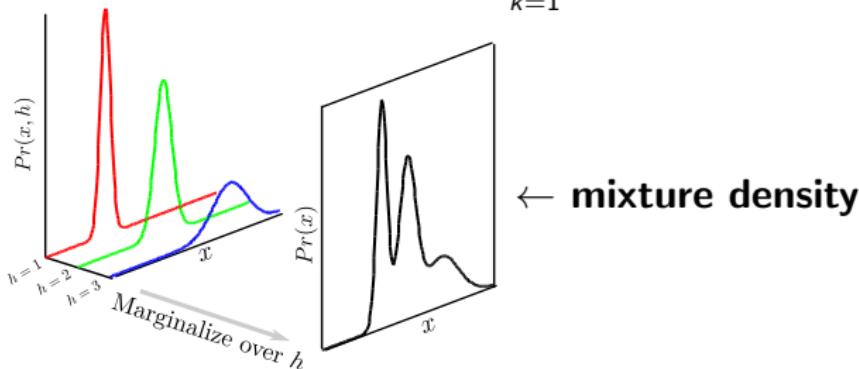
$$P(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x; \mu_k, \sigma_k^2)$$

- ▶ Learning the parameters of this model from training data x_1, \dots, x_n is not trivial - using the usual straightforward maximum likelihood approach.
- ▶ Instead learn parameters using the **Expectation-Maximization (EM)** algorithm.

Mixture of Gaussians as a marginalization

We can interpret the Mixture of Gaussians model with the introduction of a discrete hidden/latent variable h and $P(x, h)$:

$$\begin{aligned} P(x) &= \sum_{k=1}^K P(x, h = k) = \sum_{k=1}^K P(x | h = k)P(h = k) \\ &= \sum_{k=1}^K \pi_k \mathcal{N}(x; \mu_k, \sigma_k^2) \end{aligned}$$



Figures taken from **Computer Vision: models, learning and inference** by Simon Prince.

EM for two Gaussians

Assume: We know the pdf of x has this form:

$$P(x) = \pi_1 \mathcal{N}(x; \mu_1, \sigma_1^2) + \pi_2 \mathcal{N}(x; \mu_2, \sigma_2^2)$$

where $\pi_1 + \pi_2 = 1$ and $\pi_k > 0$ for components $k = 1, 2$.

Unknown: Values of the parameters (**Many!**)

$$\Theta = (\pi_1, \mu_1, \sigma_1, \mu_2, \sigma_2).$$

Have: Observed n samples x_1, \dots, x_n drawn from $P(x)$.

Want to: Estimate Θ from x_1, \dots, x_n .

How would it be possible to get them all???

EM for two Gaussians

For each sample x_i introduce a *hidden variable* h_i :

$$h_i = \begin{cases} 1 & \text{if sample } x_i \text{ was drawn from } \mathcal{N}(x; \mu_1, \sigma_1^2) \\ 2 & \text{if sample } x_i \text{ was drawn from } \mathcal{N}(x; \mu_2, \sigma_2^2) \end{cases}$$

and come up with initial values

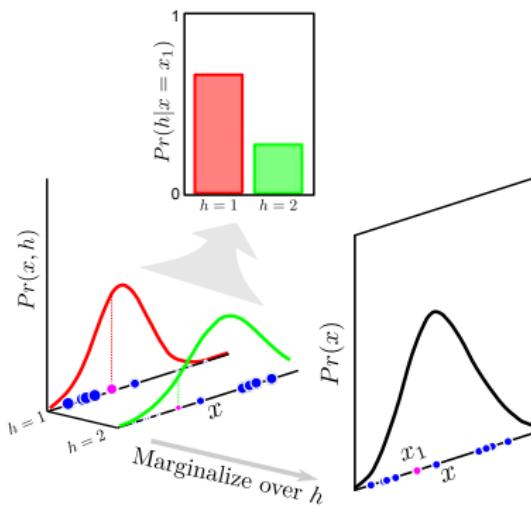
$$\Theta^{(0)} = (\pi_1^{(0)}, \mu_1^{(0)}, \sigma_1^{(0)}, \mu_2^{(0)}, \sigma_2^{(0)})$$

for each of the parameters.

EM is an *iterative algorithm* which updates $\Theta^{(t)}$ using the following two steps...

EM for two Gaussians: E-step

The responsibility of k -th Gaussian for each sample x (indicated by the size of the projected data point)



Look at each sample x along hidden variable h in the E-step

EM for two Gaussians: E-step (cont.)

E-step: Compute the “posterior probability” that x_i was generated by component k given the current estimate of the parameters $\Theta^{(t)}$. (responsibilities)

for $i = 1, \dots, n$

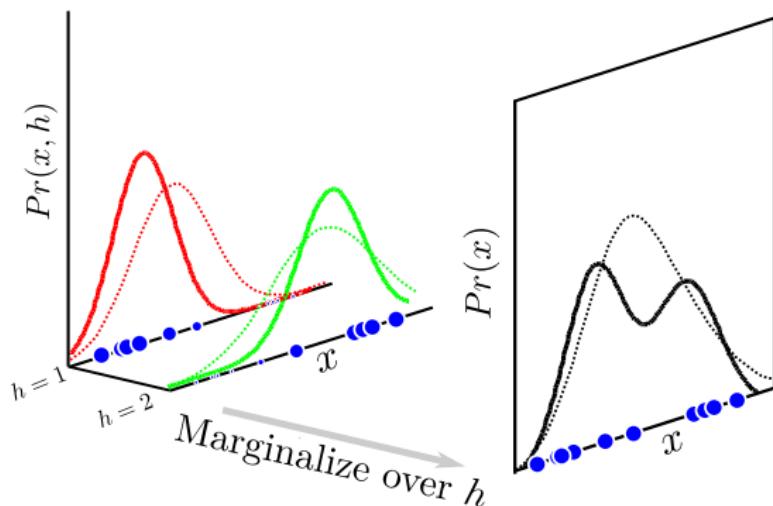
for $k = 1, 2$

$$\begin{aligned}\gamma_{ik}^{(t)} &= P(h_i = k | x_i, \Theta^{(t)}) \\ &= \frac{\pi_k^{(t)} \mathcal{N}(x_i; \mu_k^{(t)}, \sigma_k^{(t)})}{\pi_1^{(t)} \mathcal{N}(x_i; \mu_1^{(t)}, \sigma_1^{(t)}) + \pi_2^{(t)} \mathcal{N}(x_i; \mu_2^{(t)}, \sigma_2^{(t)})}\end{aligned}$$

Note: $\gamma_{i1}^{(t)} + \gamma_{i2}^{(t)} = 1$ and $\pi_1 + \pi_2 = 1$

EM for two Gaussians: M-step

Fitting the Gaussian model **for each of k -th constituent.**
Sample x_i contributes according to the responsibility γ_{ik} .



(dashed and solid lines for fit before and after update)

Look along samples x for each h in the M-step

EM for two Gaussians: M-step (cont.)

M-step: Compute the *Maximum Likelihood* of the parameters of the mixture model given out data's membership distribution, the $\gamma_i^{(t)}$'s:

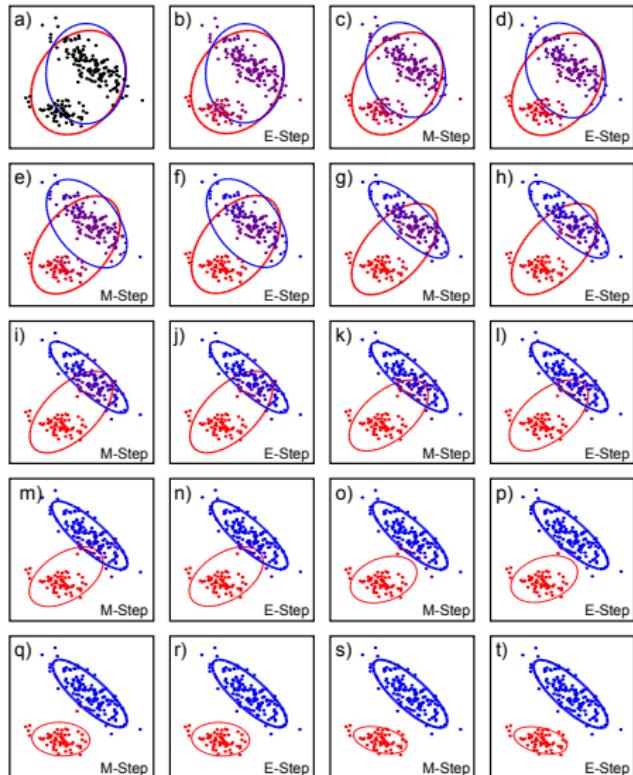
for $k = 1, 2$

$$\mu_k^{(t+1)} = \frac{\sum_{i=1}^n \gamma_{ik}^{(t)} x_i}{\sum_{i=1}^n \gamma_{ik}^{(t)}},$$

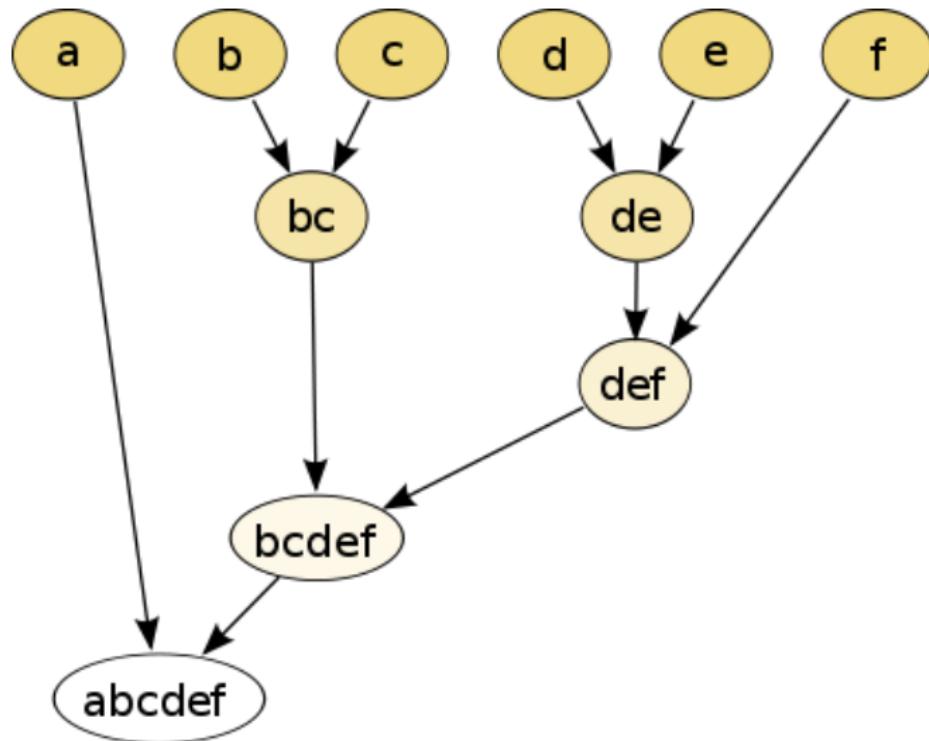
$$\sigma_k^{(t+1)} = \sqrt{\frac{\sum_{i=1}^n \gamma_{ik}^{(t)} (x_i - \mu_k^{(t+1)})^2}{\sum_{i=1}^n \gamma_{ik}^{(t)}}},$$

$$\pi_k^{(t+1)} = \frac{\sum_{i=1}^n \gamma_{ik}^{(t)}}{n}.$$

EM in practice



Hierarchical Clustering



(Figure from Wikipedia)